



Computer Systems Design

Lesson 4

Introduction to RISC-V architecture

Alexander Antonov, Assoc. Prof., ITMO University

Hangzhou, 2025

Outline the lesson

- History and motivation for development of RISC-V architecture
- Specifications
- ISA structure
- Registers view
- Instruction formats
- Base ISA, extensions

RISC-V: open RISC CPU architecture

Proposed at University of California, Berkeley (2010)

Offered multiple virtues:

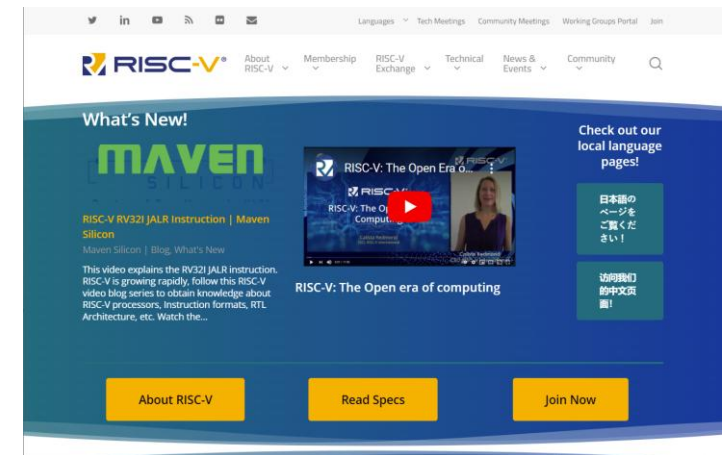
- 1) Royalty free both for academic and industrial use
- 2) High-quality, clean-slate, accumulating previous experience of HW and SW design

Published under permissive BSD license

Now run by RISC-V Foundation (Switzerland)

Official web site:

<https://riscv.org/>



Gaining big interest both from academia and industry in recent years

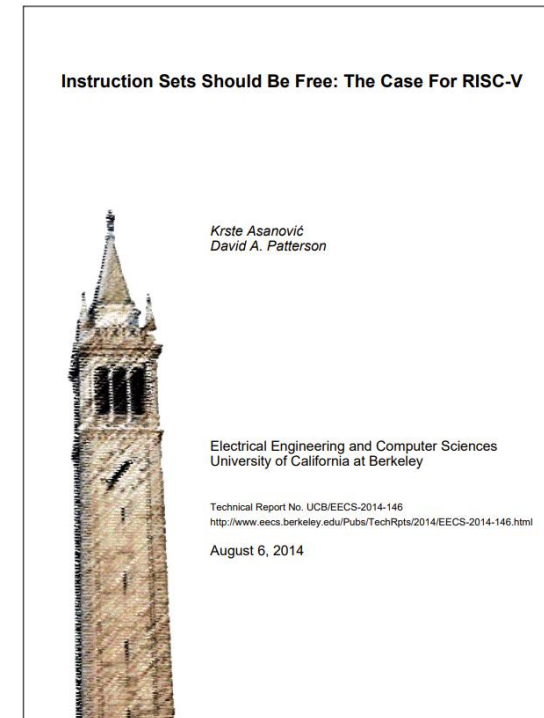
RISC-V: motivation for development

Technical Report No. UCB/EECS-2014-146:

<https://www2.eecs.berkeley.edu/Pubs/TechRpts/2014/EECS-2014-146.html>

Contains:

- 1) Justification for new ISA development
- 2) Base ISA

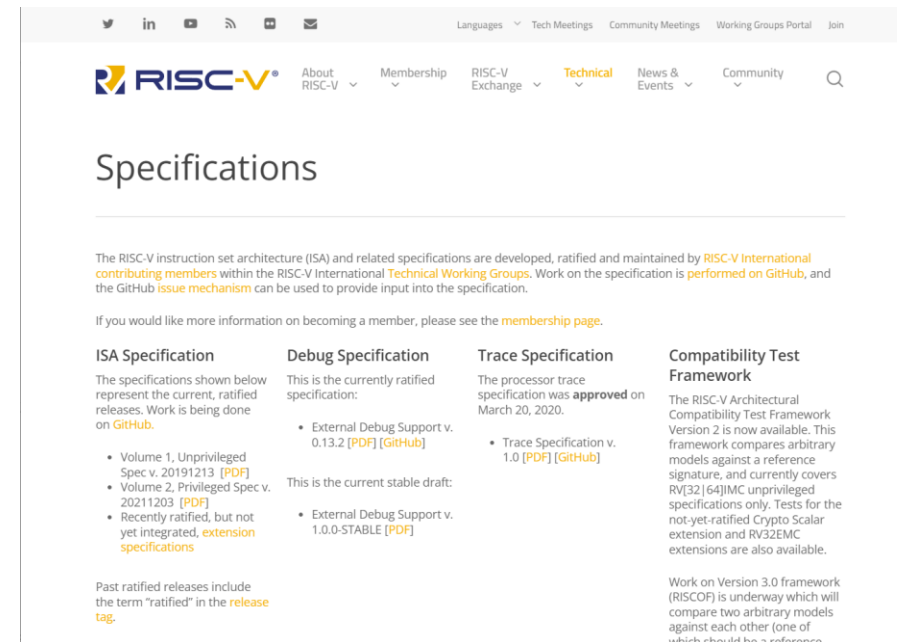


RISC-V: specifications

Available at: <https://riscv.org/technical/specifications/>

Structure of specifications:

- 1) ISA spec, Vol. 1, unprivileged spec
generic, computation-related ISA part
- 2) ISA spec, Vol. 2, privileged spec
interaction between apps, OS, hypervisors, etc.
- 3) Debug Specification
interaction with debug infrastructure
- 4) Trace Specification
interaction with trace infrastructure
- 5) Compatibility Test Framework
CPU tests proving compatibility with ISA



Used as ISA study example in our course

RISC-V: ISA structure

Designed to be modular, consists of **base (obligatory)** part and **optional extensions**

- Opcode space supports **custom extensions**
- Key elements frozen, is an active work in progress


ISA base and extensions				
Name	Description	Version	Status ^[4]	Instruction count
Base				
RVWMO	Weak Memory Ordering	2.0	Ratified	
RV32I	Base Integer Instruction Set, 32-bit	2.1	Ratified	40
RV32E	Base Integer Instruction Set (embedded), 32-bit, 16 registers	2.0	Ratified	40
RV64I	Base Integer Instruction Set, 64-bit	2.1	Ratified	15
RV64E	Base Integer Instruction Set(embedded), 64-bit	2.0	Ratified	
RV128I	Base Integer Instruction Set, 128-bit	1.7	Open	15
Extension				
M	Standard Extension for Integer Multiplication and Division	2.0	Ratified	8 (RV32) 13 (RV64)
A	Standard Extension for Atomic Instructions	2.1	Ratified	11 (RV32) 22 (RV64)
F	Standard Extension for Single-Precision Floating-Point	2.2	Ratified	26 (RV32) 30 (RV64)
D	Standard Extension for Double-Precision Floating-Point	2.2	Ratified	26 (RV32) 32 (RV64)
Zicsr	Control and Status Register (CSR) Instructions	2.0	Ratified	6
Zifencei	Instruction-Fetch Fence	2.0	Ratified	1
G	Shorthand for the IMAFDZicsr_Zifencei base and extensions	—	—	
Q	Standard Extension for Quad-Precision Floating-Point	2.2	Ratified	28 (RV32) 32 (RV64)
L	Standard Extension for Decimal Floating-Point	0.0	Open	
C	Standard Extension for Compressed Instructions	2.0	Ratified	40
B	Standard Extension for Bit Manipulation	1.0	Ratified	43 ^[40]
J	Standard Extension for Dynamically Translated Languages	0.0	Open	
T	Standard Extension for Transactional Memory	0.0	Open	
P	Standard Extension for Packed-SIMD Instructions	0.9.10	Open	
V	Standard Extension for Vector Operations	1.0	Frozen	187 ^[41]
Zk	Standard Extension for Scalar Cryptography	1.0.1	Ratified	49 ^[42]
H	Standard Extension for Hypervisor	1.0	Ratified	15
S	Standard Extension for Supervisor-level Instructions	1.12	Ratified	4
Zam	Misaligned Atomics	0.1	Open	
Zhintpause	Pause Hint	2.0	Ratified	
Zhintntli	Non-Temporal Locality Hints	0.2	Open	
Zfa	Additional Floating-Point Instructions	0.1	Open	
Zfh	Half-Precision Floating-Point	1.0	Ratified	
Zfhmin	Minimal Half-Precision Floating-Point	1.0	Ratified	
Zfmx	Single-Precision Floating-Point in Integer Register	1.0	Ratified	
Zdinx	Double-Precision Floating-Point in Integer Register	1.0	Ratified	
Zhinx	Half-Precision Floating-Point in Integer Register	1.0	Ratified	
Zhinxmin	Minimal Half-Precision Floating-Point in Integer Register	1.0	Ratified	
Zmmul	Multiplication Subset of the M Extension	1.0	Ratified	
Ztso	Total Store Ordering	1.0	Ratified	

Base ISA: register space (RV32I)

Register bank consists of 32x **32-bit general-purpose registers** (XLEN=32) + program counter

Only x0 register is special – always zero (writes don't take effect)

(ISA spec, Vol. 1, unprivileged spec, p. 14)




Base ISA: register space (RV32E)

Instead of 32x, 16x 32-bit registers are provided

Targeted for embedded applications

(ISA spec, Vol. 1, unprivileged spec, p. 33)



Base ISA: instruction formats

Base instructions are **32-bit data words** assembled according to the following formats (ISA spec, Vol. 1, unprivileged spec, p. 16):

Priorities:

- Unification of instructions sizes and formats (preserving high code density)
- Decreasing diversity of possible locations of instruction fields for different types of instructions

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0		
funct7				rs2		rs1		funct3		rd			opcode		R-type	
imm[11:0]						rs1		funct3		rd			opcode		I-type	
imm[11:5]				rs2		rs1		funct3		imm[4:0]			opcode		S-type	
imm[12]		imm[10:5]		rs2		rs1		funct3		imm[4:1]		imm[11]		opcode		B-type
imm[31:12]										rd			opcode		U-type	
imm[20]		imm[10:1]			imm[11]		imm[19:12]			rd			opcode		J-type	

Base RV32I ISA

RV32I Base Instruction Set							
imm[31:12]				rd		0110111	LUI
imm[31:12]				rd		0010111	AUIPC
imm[20:10:11:19:12]				rd		1101111	JAL
imm[11:0]				rd		1100111	JALR
imm[12:10:5]	rs2	rs1	000	imm[4:1:11]	1100011	BEQ	
imm[12:10:5]	rs2	rs1	001	imm[4:1:11]	1100011	BNE	
imm[12:10:5]	rs2	rs1	100	imm[4:1:11]	1100011	BLT	
imm[12:10:5]	rs2	rs1	101	imm[4:1:11]	1100011	BGE	
imm[12:10:5]	rs2	rs1	110	imm[4:1:11]	1100011	BLTU	
imm[12:10:5]	rs2	rs1	111	imm[4:1:11]	1100011	BGEU	
imm[11:0]				rd		0000011	LB
imm[11:0]				rd		0000011	LH
imm[11:0]				rd		0000011	LW
imm[11:0]				rd		0000011	LBU
imm[11:0]				rd		0000011	LHU
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011	SB	
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011	SH	
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	SW	
imm[11:0]				rd		0010011	ADDI
imm[11:0]				rd		0010011	SLTI
imm[11:0]				rd		0010011	SLTIU
imm[11:0]				rd		0010011	XORI
imm[11:0]				rd		0010011	ORI
imm[11:0]				rd		0010011	ANDI
0000000		shamt	rs1	001	rd	0010011	SLLI
0000000		shamt	rs1	101	rd	0010011	SRLI
0100000		shamt	rs1	101	rd	0010011	SRAI
0000000		rs2	rs1	000	rd	0110011	ADD
0100000		rs2	rs1	000	rd	0110011	SUB
0000000		rs2	rs1	001	rd	0110011	SLL
0000000		rs2	rs1	010	rd	0110011	SLT
0000000		rs2	rs1	011	rd	0110011	SLTU
0000000		rs2	rs1	100	rd	0110011	XOR
0000000		rs2	rs1	101	rd	0110011	SRL
0100000		rs2	rs1	101	rd	0110011	SRA
0000000		rs2	rs1	110	rd	0110011	OR
0000000		rs2	rs1	111	rd	0110011	AND
fin	pred	succ	rs1	000	rd	0001111	FENCE
000000000000			00000	000	00000	1110011	ECALL
000000000001			00000	000	00000	1110011	EBREAK

Write constant to register (high bits)

Writing PC+offset to register
(for pc-relative address construction)

Branches with memorizing return points
(for procedure calls)

Conditional branches

Load/store instructions
(transfers between registers and
main memory)

Arithmetic and logical operations

Special-purpose instructions

M extension

RV32M Standard Extension						
0000001	rs2	rs1	000	rd	0110011	MUL
0000001	rs2	rs1	001	rd	0110011	MULH
0000001	rs2	rs1	010	rd	0110011	MULHSU
0000001	rs2	rs1	011	rd	0110011	MULHU
0000001	rs2	rs1	100	rd	0110011	DIV
0000001	rs2	rs1	101	rd	0110011	DIVU
0000001	rs2	rs1	110	rd	0110011	REM
0000001	rs2	rs1	111	rd	0110011	REMU

Multiplication

Division

Remainder of the division

C extension

Shorter (16-bit) instructions that can replace ~50% traditional instructions
(resulting in 25–30% total code size reduction)

Format	Meaning	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
CR	Register	funct4				rd/rs1				rs2				op					
CI	Immediate	funct3		imm		rd/rs1				imm				op					
CSS	Stack-relative Store	funct3		imm						rs2				op					
CIW	Wide Immediate	funct3		imm								rd'		op					
CL	Load	funct3		imm			rs1'			imm		rd'		op					
CS	Store	funct3		imm			rs1'			imm		rs2'		op					
CA	Arithmetic	funct6					rd'/rs1'			funct2		rs2'		op					
CB	Branch	funct3		offset			rs1'			offset				op					
CJ	Jump	funct3		jump target												op			

RVC Register Number	000	001	010	011	100	101	110	111
Integer Register Number	x8	x9	x10	x11	x12	x13	x14	x15
Integer Register ABI Name	s0	s1	a0	a1	a2	a3	a4	a5
Floating-Point Register Number	f8	f9	f10	f11	f12	f13	f14	f15
Floating-Point Register ABI Name	fs0	fs1	fa0	fa1	fa2	fa3	fa4	fa5

ISA spec, Vol. 1, unprivileged spec, p. 100

A extension

Atomic read/modify/write operation to *avoid race conditions* between multiple threads/core

RV32A Standard Extension

00010	aq	rl	00000	rs1	010	rd	0101111	LR.W
00011	aq	rl	rs2	rs1	010	rd	0101111	SC.W
00001	aq	rl	rs2	rs1	010	rd	0101111	AMOSWAP.W
00000	aq	rl	rs2	rs1	010	rd	0101111	AMOADD.W
00100	aq	rl	rs2	rs1	010	rd	0101111	AMOXOR.W
01100	aq	rl	rs2	rs1	010	rd	0101111	AMOAND.W
01000	aq	rl	rs2	rs1	010	rd	0101111	AMOOR.W
10000	aq	rl	rs2	rs1	010	rd	0101111	AMOMIN.W
10100	aq	rl	rs2	rs1	010	rd	0101111	AMOMAX.W
11000	aq	rl	rs2	rs1	010	rd	0101111	AMOMINU.W
11100	aq	rl	rs2	rs1	010	rd	0101111	AMOMAXU.W

“Load-reserved” – load and reserve exclusive rights to the word

“Store-conditional” – store is having rights

Atomic read/modify/write operations

-
- Diagram illustrating the IEEE 754 single-precision floating-point format. The 32-bit word is divided into three fields:
- sign** (1 bit): Bit 31.
 - exponent** (8 bits): Bits 30 to 23.
 - fraction** (23 bits): Bits 22 to 0.
- The bit pattern shown is: $0 \ 01111000 \ 01000000000000000000000$. The decimal value represented is 0.15625 .

RV32D Standard Extension							
imm[11:0]			rs1	011	rd	0000111	FLD
imm[11:5]		rs2	rs1	011	imm[4:0]	0100111	FSD
rs3	01	rs2	rs1	rm	rd	1000011	FMADD.D
rs3	01	rs2	rs1	rm	rd	1000111	FMSUB.D
rs3	01	rs2	rs1	rm	rd	1001011	FNMSUB.D
rs3	01	rs2	rs1	rm	rd	1001111	FNMADD.D
0000001		rs2	rs1	rm	rd	1010011	FADD.D
0000101		rs2	rs1	rm	rd	1010011	FSUB.D
0001001		rs2	rs1	rm	rd	1010011	FMUL.D
0001101		rs2	rs1	rm	rd	1010011	FDIV.D
0101101		00000	rs1	rm	rd	1010011	FSQRT.D
0010001		rs2	rs1	000	rd	1010011	FSGNJ.D
0010001		rs2	rs1	001	rd	1010011	FSGNJN.D
0010001		rs2	rs1	010	rd	1010011	FSGNJX.D
0010101		rs2	rs1	000	rd	1010011	FMIN.D
0010101		rs2	rs1	001	rd	1010011	FMAX.D
0100000		00001	rs1	rm	rd	1010011	FCVT.S.D
0100001		00000	rs1	rm	rd	1010011	FCVT.D.S
1010001		rs2	rs1	010	rd	1010011	FEQ.D
1010001		rs2	rs1	001	rd	1010011	FLT.D
1010001		rs2	rs1	000	rd	1010011	FLED
1110001		00000	rs1	001	rd	1010011	FCLASS.D
1100001		00000	rs1	rm	rd	1010011	FCVT.W.D
1100001		00001	rs1	rm	rd	1010011	FCVT.W.U.D
1101001		00000	rs1	rm	rd	1010011	FCVT.D.W
1101001		00001	rs1	rm	rd	1010011	FCVT.D.W.U

Zicsr extension

31	20	19	15	14	12	11	7	6	0	
csr			rs1		funct3		rd		opcode	
12			5		3		5		7	
source/dest			source		CSRRW		dest		SYSTEM	
source/dest			source		CSRRS		dest		SYSTEM	
source/dest			source		CSRRC		dest		SYSTEM	
source/dest			uimm[4:0]		CSRRWI		dest		SYSTEM	
source/dest			uimm[4:0]		CSRRSI		dest		SYSTEM	
source/dest			uimm[4:0]		CSRRCI		dest		SYSTEM	

Binary and mnemonic representation of instructions

Mnemonic of instructions has unique binary representation (and vice versa in most cases)

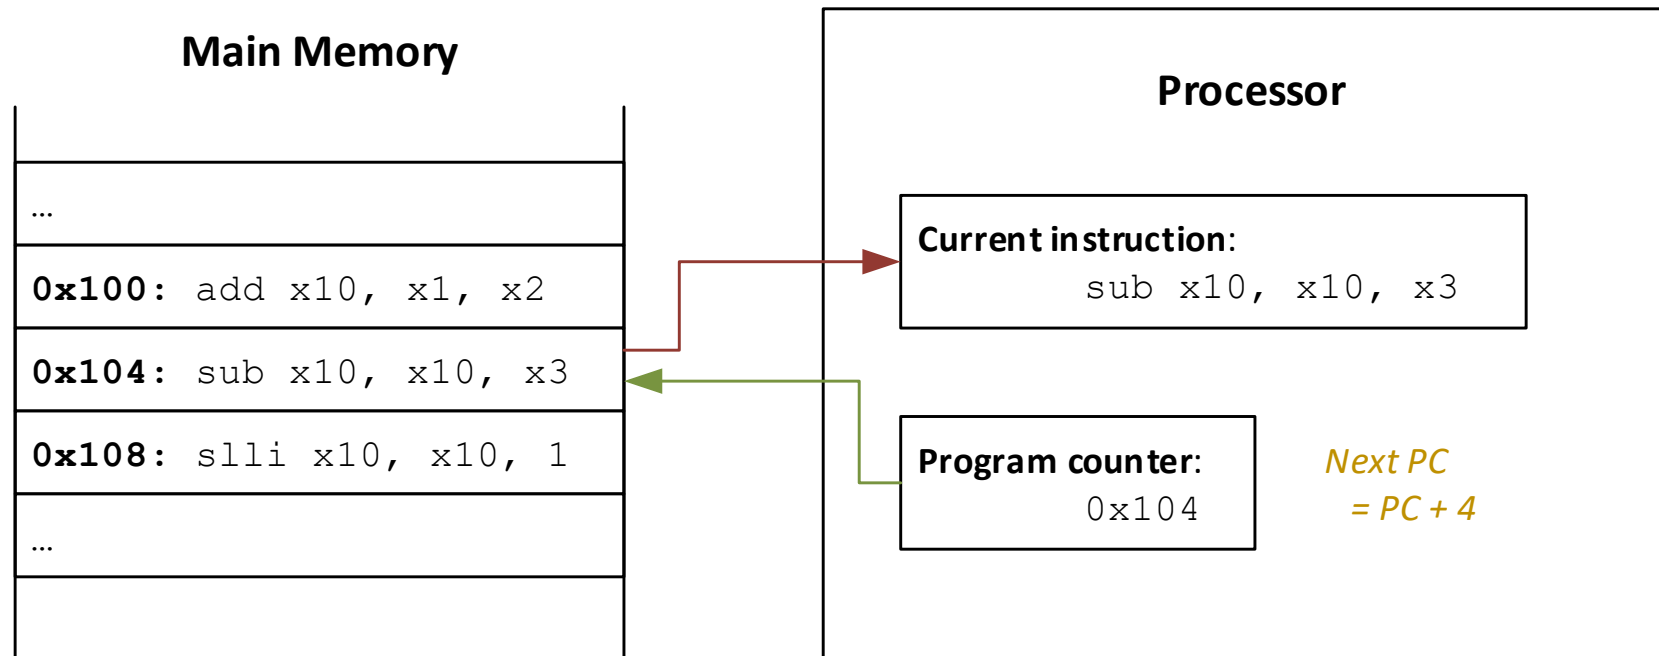
Mnemonic -> binary:

- 1) Find instruction type and format
- 2) Fill in the instruction fields

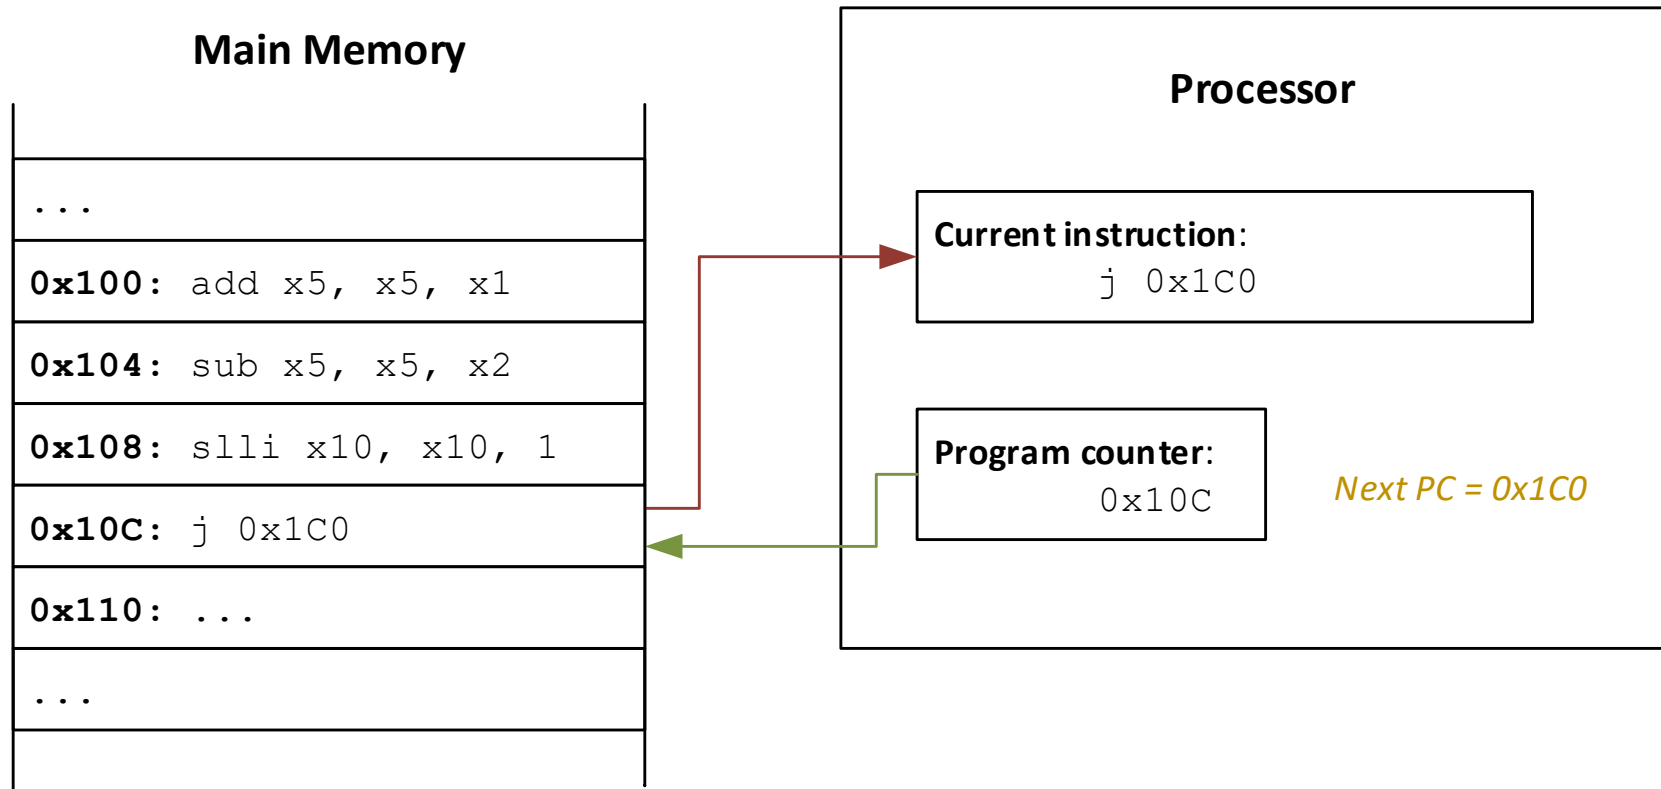
Binary -> mnemonic:

- 1) Separate instruction opcode
- 2) Reveal instruction type
- 3) Separate individual fields

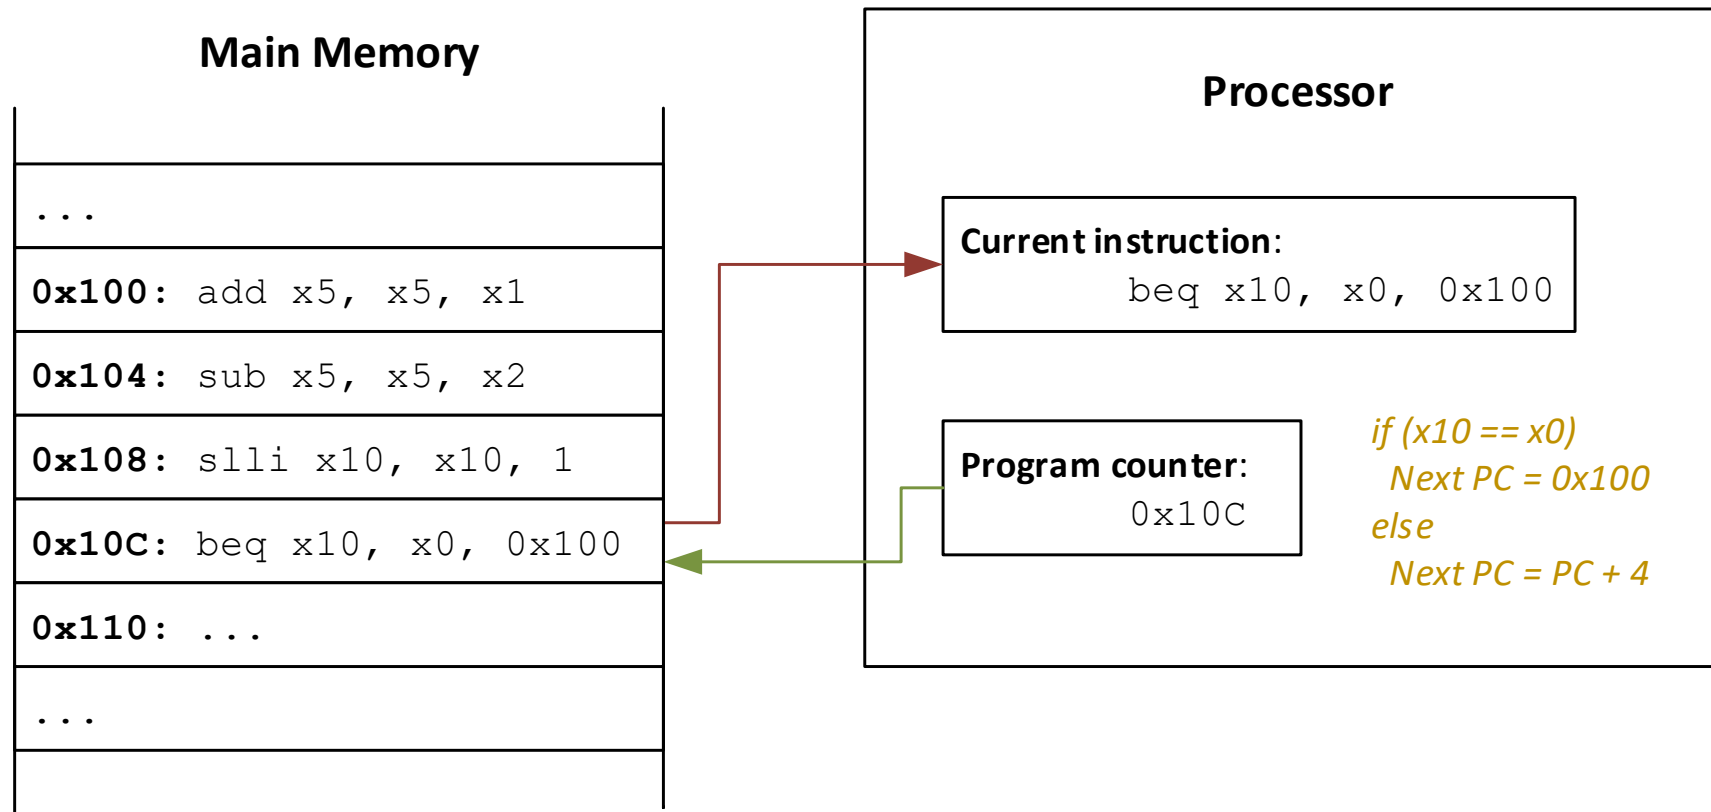
Example: linear code block



Example: unconditional branch (infinite loop)



Example: conditional branch





Thank you for the lesson!

Alexander Antonov, Assoc. Prof., antonov@itmo.ru

Hangzhou, 2025