



Computer Systems Design

Lesson 7

Basics of Assembly programming

Alexander Antonov, Assoc. Prof., ITMO University

Hangzhou, 2025

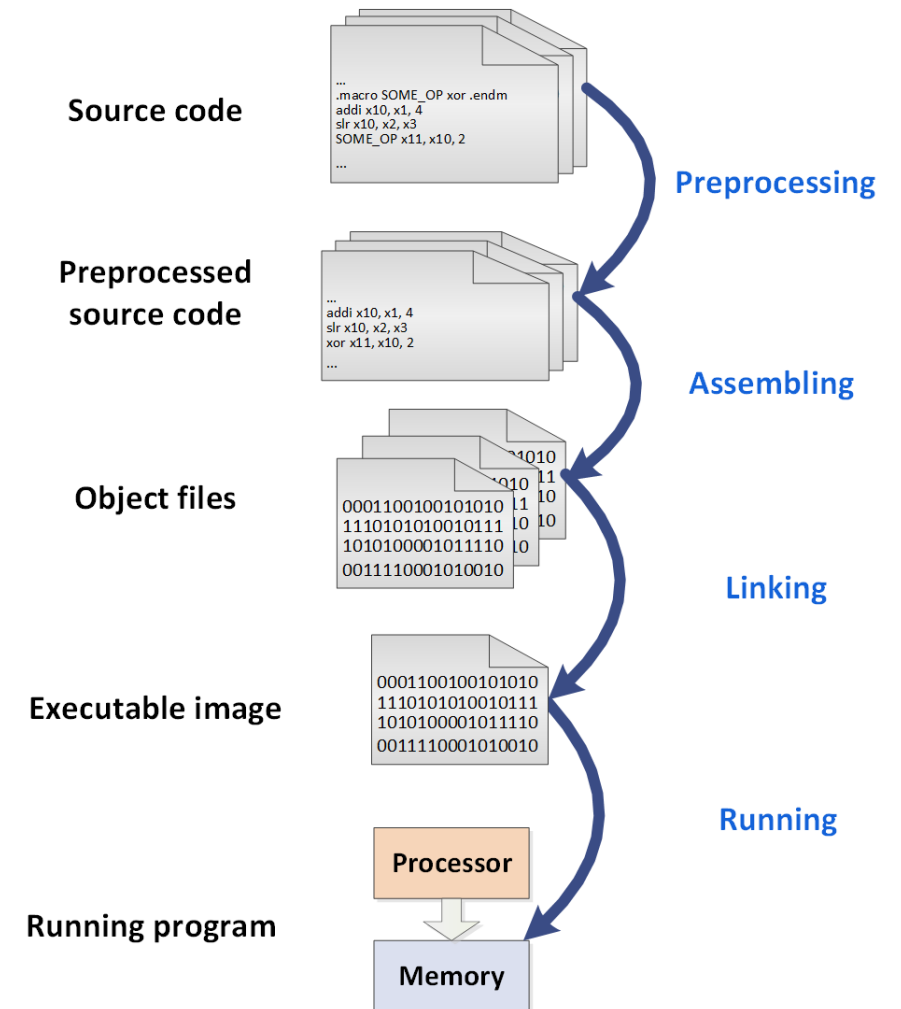
Outline the lesson

- Instructions and pseudo-instructions
- C vs. ASM idioms
 - constants and variables
 - branching, loops
 - procedures
 - callbacks
- Inline assembly

Assembly language

Assembler languages: languages composed of mnemonic (textual) representations of processor instructions

- Specific assembler is tied to certain processor ISA
- Single assembler command represents one or several processor instructions
- No complex optimizations are implemented
- All binary code should be unambiguously understood (“decoded”) by the processor



Pseudo-instructions

Pseudo-instructions: assembly instructions without exclusive correspondence to machine instructions
Simplify manual assembly programming

| | | |
|---------------------|-------------------------|---------------------------------|
| nop | addi x0, x0, 0 | No operation |
| li rd, immediate | <i>Myriad sequences</i> | Load immediate |
| mv rd, rs | addi rd, rs, 0 | Copy register |
| not rd, rs | xori rd, rs, -1 | One's complement |
| neg rd, rs | sub rd, x0, rs | Two's complement |
| negw rd, rs | subw rd, x0, rs | Two's complement word |
| sext.w rd, rs | addiw rd, rs, 0 | Sign extend word |
| seqz rd, rs | sltiu rd, rs, 1 | Set if = zero |
| snez rd, rs | sltu rd, x0, rs | Set if ≠ zero |
| sltz rd, rs | slt rd, rs, x0 | Set if < zero |
| sgtz rd, rs | slt rd, x0, rs | Set if > zero |
| fmv.s rd, rs | fsgnj.s rd, rs, rs | Copy single-precision register |
| fabs.s rd, rs | fsgnjx.s rd, rs, rs | Single-precision absolute value |
| fneg.s rd, rs | fsgnjn.s rd, rs, rs | Single-precision negate |
| fmv.d rd, rs | fsgnj.d rd, rs, rs | Copy double-precision register |
| fabs.d rd, rs | fsgnjx.d rd, rs, rs | Double-precision absolute value |
| fneg.d rd, rs | fsgnjd.d rd, rs, rs | Double-precision negate |
| beqz rs, offset | beq rs, x0, offset | Branch if = zero |
| bnez rs, offset | bne rs, x0, offset | Branch if ≠ zero |
| blez rs, offset | bge x0, rs, offset | Branch if ≤ zero |
| bgez rs, offset | bge rs, x0, offset | Branch if ≥ zero |
| bltz rs, offset | blt rs, x0, offset | Branch if < zero |
| bgtz rs, offset | blt x0, rs, offset | Branch if > zero |
| bgt rs, rt, offset | blt rt, rs, offset | Branch if > |
| ble rs, rt, offset | bge rt, rs, offset | Branch if ≤ |
| bgtu rs, rt, offset | bltu rt, rs, offset | Branch if >, unsigned |
| bleu rs, rt, offset | bgeu rt, rs, offset | Branch if ≤, unsigned |

Table 25.2: RISC-V pseudoinstructions.

| pseudoinstruction | Base Instruction | Meaning |
|-------------------|--|-----------------------------------|
| j offset | jal x0, offset | Jump |
| jal offset | jal x1, offset | Jump and link |
| jr rs | jalr x0, 0(rs) | Jump register |
| jalr rs | jalr x1, 0(rs) | Jump and link register |
| ret | jalr x0, 0(x1) | Return from subroutine |
| call offset | auipc x1, offset[31:12] + offset[11] jalr x1, offset[11:0] (x1) | Call far-away subroutine |
| tail offset | auipc x6, offset[31:12] + offset[11] jalr x0, offset[11:0] (x6) | Tail call far-away subroutine |
| fence | fence iorw, iorw | Fence on all memory and I/O |
| rdinstret[h] rd | csrrs rd, instret[h], x0 | Read instructions-retired counter |
| rdcycle[h] rd | csrrs rd, cycle[h], x0 | Read cycle counter |
| rdtime[h] rd | csrrs rd, time[h], x0 | Read real-time clock |
| csrr rd, csr | csrrs rd, csr, x0 | Read CSR |
| csrw csr, rs | csrrw x0, csr, rs | Write CSR |
| csrs csr, rs | csrrs x0, csr, rs | Set bits in CSR |
| csrc csr, rs | csrrc x0, csr, rs | Clear bits in CSR |
| csrwi csr, imm | csrrwi x0, csr, imm | Write CSR, immediate |
| csrsi csr, imm | csrrsi x0, csr, imm | Set bits in CSR, immediate |
| csrci csr, imm | csrrci x0, csr, imm | Clear bits in CSR, immediate |
| frcsr rd | csrrs rd, fcsr, x0 | Read FP control/status register |
| fscsr rd, rs | csrrw rd, fcsr, rs | Swap FP control/status register |
| fscsr rs | csrrw x0, fcsr, rs | Write FP control/status register |
| frfm rd | csrrs rd, frm, x0 | Read FP rounding mode |
| fsrm rd, rs | csrrw rd, frm, rs | Swap FP rounding mode |
| fsrm rs | csrrw x0, frm, rs | Write FP rounding mode |
| frflags rd | csrrs rd, fflags, x0 | Read FP exception flags |
| fsflags rd, rs | csrrw rd, fflags, rs | Swap FP exception flags |
| fsflags rs | csrrw x0, fflags, rs | Write FP exception flags |

Table 25.3: RISC-V pseudoinstructions.

ISA spec, Vol. 1, unprivileged spec, Chapter 25

C vs. ASM: constant initialization

“Short” constant

```
unsigned int some_var = 134;  
li a4,134    [addi a4, x0, 134]
```

“Long” constant (longer than 12-bit imm in I-type instruction)

```
unsigned int some_var = 3489110677; // 0xcff79a95  
lui    a5,0xcff7a    # constructing 0xcff79... part  
addi   a5,a5,-1387   # constructing 0x.....a95 part
```

C vs. ASM: working with main memory

```
// memory-mapped registers at addresses 0x80000000 and 0x80000004
#define IO_IN      (*(volatile unsigned int *)(0x80000000))
#define IO_OUT     (*(volatile unsigned int *)(0x80000004))

IO_OUT = IO_IN + 211;
```

```
lui a4,0x80000    # writing high part of IO_* address to a4
lw  a5,0(a4)      # loading IO_IN memory cell to a5

addi a5,a5,211    # adding 211
sw  a5,4(a4)      # storing a5 to IO_OUT memory cell
```

C vs. ASM: branching

```
// memory-mapped registers at addresses 0x80000000 and 0x80000004
#define IO_IN      (*(volatile unsigned int *) (0x80000000))
#define IO_OUT     (*(volatile unsigned int *) (0x80000004))

if (IO_IN == 15) {
    IO_OUT = IO_OUT + 10;
} else {
    IO_OUT = IO_OUT + 20;
}
while(1) {} // infinite loop
```

```
518:    lui    a5,0x80000
51c:    lw     a3,0(a5)      # loading IO_IN to a3
520:    li     a4,15
524:    beq    a3,a4,538     # branch if IO_IN == 15
528:    lw     a4,4(a5)
52c:    addi   a4,a4,20
530:    sw     a4,4(a5)      # storing a4 to IO_OUT
534:    j      534           # infinite loop
538:    lw     a4,4(a5)
53c:    addi   a4,a4,10
540:    sw     a4,4(a5)      # storing a4 to IO_OUT
544:    j      534           # go to infinite loop
```

C vs. ASM: arrays and loops

```
#define ARR_SIZE 256

int array [ARR_SIZE];

for (int i=0; i<ARR_SIZE; i++) {
    array[i] = i;          // initializing array with counter values
}
while(1) {} // infinite loop
```

| | | | |
|------|------|----------------------|--|
| 51c: | li | a5,0 | # writing index of element to a5 |
| 520: | li | a3,256 | # writing array size to a3 |
| 524: | slli | a4,a5,0x2 | # writing element address offset to a4 |
| 528: | addi | a2,sp,<array addr> | # writing array base address to a2 |
| 52c: | add | a4,a2,a4 | # writing element address to a4 |
| 530: | sw | a5,0(a4) | # writing element data |
| 534: | addi | a5,a5,1 | # incrementing index |
| 538: | bne | a5,a3,524 <main+0xc> | # checking exit condition |
| 53c: | j | 53c <main+0x24> | # infinite loop |

Procedures

Procedure: reusable block of code than can be called from multiple locations

Typically has *input arguments* and *output value*

```
int FindMax(int src1, int src2) {           // procedure (callee)
    if (src1 > src2) return src1;
    else return src2;
}

int DoSmth1()                               // parent procedure (caller 1)
{
    IO_OUT = FindMax(1, 2);                 // procedure call
    IO_OUT = FindMax(24, 7);                // procedure call
    IO_OUT = FindMax(19, 138);              // procedure call
    // some other code ...
}

int DoSmth2()                               // parent procedure (caller 2)
{
    IO_OUT = FindMax(14, 15);               // procedure call
    // some other code ...
}
```

Recall: Application Binary Interface (ABI)

Convention of register usage by programs

Defines *recommended* role of general-purpose registers for procedure calls, stacks, etc.

Followed by programmers and compiler developers

ISA spec, Vol. 1, unprivileged spec, p. 137

| Register | ABI Name | Description | Saver |
|----------|----------|-----------------------------------|--------|
| x0 | zero | Hard-wired zero | — |
| x1 | ra | Return address | Caller |
| x2 | sp | Stack pointer | Callee |
| x3 | gp | Global pointer | — |
| x4 | tp | Thread pointer | — |
| x5 | t0 | Temporary/alternate link register | Caller |
| x6–7 | t1–2 | Temporaries | Caller |
| x8 | s0/fp | Saved register/frame pointer | Callee |
| x9 | s1 | Saved register | Callee |
| x10–11 | a0–1 | Function arguments/return values | Caller |
| x12–17 | a2–7 | Function arguments | Caller |
| x18–27 | s2–11 | Saved registers | Callee |
| x28–31 | t3–6 | Temporaries | Caller |
| f0–7 | ft0–7 | FP temporaries | Caller |
| f8–9 | fs0–1 | FP saved registers | Callee |
| f10–11 | fa0–1 | FP arguments/return values | Caller |
| f12–17 | fa2–7 | FP arguments | Caller |
| f18–27 | fs2–11 | FP saved registers | Callee |
| f28–31 | ft8–11 | FP temporaries | Caller |

“Jump and Link” instructions

To return from functions *right after the calling point*, this location has to be saved

“Jump and link” instructions allow to save PC+4 in register (to know where to return when procedure finished)

Jump and link (JAL)

```
jal rd, <procedure offset>
```

Performs:

- $rd \leftarrow PC + 4$
- $PC \leftarrow PC + \text{procedure offset (20 bits)}$

Jump and link register (JALR)

```
jal rd, rs, <procedure offset>
```

Performs:

- $rd \leftarrow PC + 4$
- $PC \leftarrow rs + \text{procedure offset (12 bits)}$

C vs. ASM: procedures

```
#define IO_IN0      (*(volatile unsigned int *) (0x80000000))
#define IO_IN1      (*(volatile unsigned int *) (0x80000004))
#define IO_OUT      (*(volatile unsigned int *) (0x80000008))
```

```
int Compute(int x0, int x1) {           // procedure (callee)
    if (x0 > 10) return (x0 >> 3) | (x1 + 456);
    else return (x0 << 2) & (x1 - 211);
}
```

```
int main()                             // procedure (caller)
{
    IO_OUT = Compute(IO_IN0, IO_IN1);
    while(1) {} // infinite loop
}
```

```
000002d8 <Compute>: ## procedure: callee
2d8: 00a00793      li      a5,10
2dc: 00a7da63      bge     a5,a0,2f0 <Compute+0x18>    # branching depending on (x0 > 10)
2e0: 1c858593      addi    a1,a1,456
2e4: 40355513      srai    a0,a0,0x3
2e8: 00b56533      or      a0,a0,a1
2ec: 00008067      ret
2f0: f2d58593      addi    a1,a1,-211
2f4: 00251513      slli    a0,a0,0x2
2f8: 00b57533      and     a0,a0,a1
2fc: 00008067      ret
# returning to caller

00000540 <main>:   ## procedure: caller
540: 80000437      lui     s0,0x800000    # saving IO_* base address to s0
544: 00442503      lw      a0,0(s0)       # saving argument 0 to a0
548: 00042583      lw      a1,4(s0)       # saving argument 1 to a1
54c: 00112623      sw      ra,12(sp)      # saving return address
550: d81ff0ef      jal     ra,2d8 <Compute> # calling procedure
554: 00a42023      sw      a0,8(s0)       # writing return value from a0 to IO_OUT
558: 0000006f      j       558 <main+0x20> # infinite loop
```

C vs. ASM: callbacks (C)

Sometimes it is useful to pass *references to functions*

```
int call_function_by_reference(int (call_vector)(void)) {  
    return call_vector();  
}  
  
int func0() {  
    return IO_IN0 + 3;  
}  
  
int func1() {  
    return IO_IN1 + 5;  
}  
  
int main()  
{  
    if (IO_IN0 == 0) IO_OUT = call_function_by_reference(&func0);  
    else IO_OUT = call_function_by_reference(&func1);  
  
    while(1) {}           // infinite loop  
}
```

C vs. ASM: callbacks (ASM)

```

000002a4 <func0>:
2a4: 800007b7      lui    a5,0x800000
2a8: 0007a503      lw     a0,0(a5)          # loading value at address 0x80000000 to a0
2ac: 00350513      addi   a0,a0,3
2b0: 00008067      ret

000002b4 <func1>:
2b4: 800007b7      lui    a5,0x800000
2b8: 0047a503      lw     a0,4(a5)          # loading value at address 0x80000004 to a0
2bc: 00550513      addi   a0,a0,5
2c0: 00008067      ret

000002f8 <call_function_by_reference>:
2f8: 00050313      mv     t1,a0
2fc: 00030067      jr     t1                # calling passed function reference

00000540 <main>:
540: ff010113      addi   sp,sp,-16
544: 00812423      sw     s0,8(sp)
548: 80000437      lui    s0,0x800000
54c: 00042783      lw     a5,0(s0)          # loading value at address 0x80000000 to a5
550: 00112623      sw     ra,12(sp)
554: 00079c63      bnez   a5,56c
558: 00000517      auipc  a0,0x0
55c: d4c50513      addi   a0,a0,-692         # loading reference 0x2a4 <func0> to a0
560: d99ff0ef      jal    ra,2f8             # calling <call_function_by_reference>
564: 00a42423      sw     a0,8(s0)          # storing result to IO_OUT
568: 0000006f      j      568               # infinite loop
56c: 00000517      auipc  a0,0x0
570: d4850513      addi   a0,a0,-696         # loading reference 0x2b4 <func1> to a0
574: d85ff0ef      jal    ra,2f8             # calling <call_function_by_reference>
578: 00a42423      sw     a0,8(s0)          # storing result to IO_OUT
57c: fedff06f      j      568               # going to infinite loop

```

Inline assembly in C

C and assembly execution models are very close to each other

C and assembly entry can be ***mixed*** in programs

```
#define IO_IN0      (*(volatile unsigned int *)(0x80000000))
#define IO_IN1      (*(volatile unsigned int *)(0x80000004))
#define IO_OUT      (*(volatile unsigned int *)(0x80000008))
```

```
int main() {
    //IO_OUT = Compute(IO_IN, IO_OUT);
```

```
    int x0 = IO_IN0;
    int x1 = IO_IN1;
    int y  = 0;

    asm volatile (
        " addi %2, %0, 5 \n\t"
        " xor %2, %2, %1 \n\t"
        " andi %2, %2, 288 \n\t"
        : "=r" (y)
        : "r" (x0), "r" (x1));
```

```
    IO_OUT = y;
```

```
    while(1) {} // infinite loop
```

```
}
```

```
00000518 <main>:
518:      lui   a5,0x80000
51c:      lw    a4,0(a5)          # 80000000
520:      lw    a3,4(a5)
524:      addi  a3,a4,5
528:      xor   a3,a3,a4
52c:      andi  a3,a3,288
530:      sw    a4,8(a5)
534:      j     534 <main+0x1c>
```



Thank you for the lesson!

Alexander Antonov, Assoc. Prof., antonov@itmo.ru

Hangzhou, 2025