



# Computer Systems Design

## Lesson 11

### Control flow optimization

Alexander Antonov, Assoc. Prof., ITMO University

Hangzhou, 2025

## Outline the lesson

- Generic strategies of data directed control flow implementation
- Control flow implementations in pure hardware
- Control flow implementations in processor microarchitecture
- Control flow optimizations in compilers
- Control flow optimizations in programming

# Control and Data Flow Graph

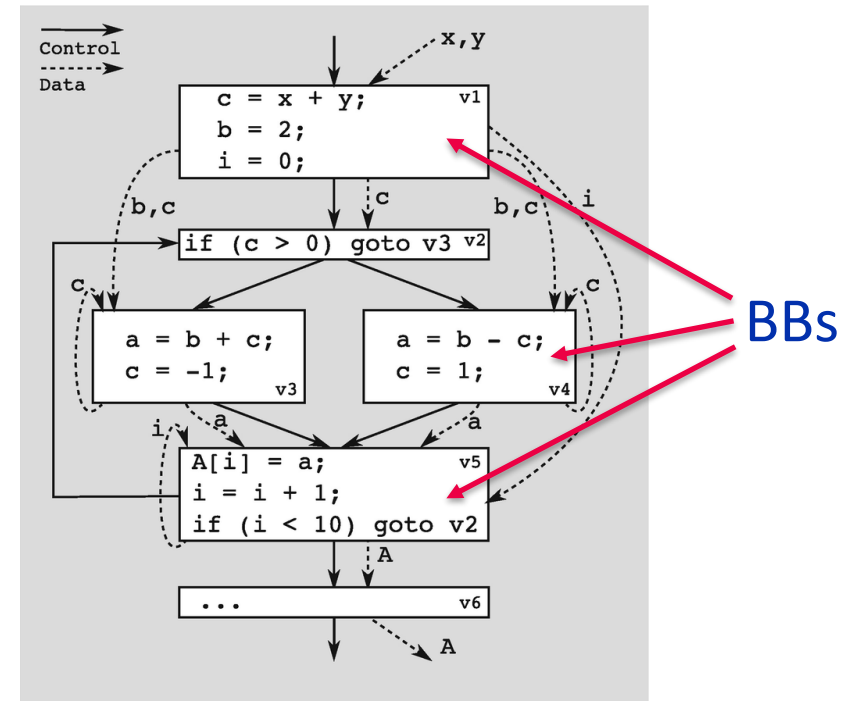
CDFG (Control and Data Flow Graph) – common IR for computer programs in compilers

**Basic blocks (BBs)** have one entry and one exit point

Optimizations are split in applied:

- **within BBs** (data flow optimizations)
- **between BBs** (control flow optimizations)

```
1: ...  
2: c = x + y;  
3: b = 2;  
4: for (i = 0; i < 10; i++) {  
5:     if (c > 0) {  
6:         a = b + c;  
7:         c = -1;  
8:     } else {  
9:         a = b - c;  
10:        c = 1;  
11:    }  
12:    A[i] = a;  
13: }  
14: ...
```




## Data-directed control flow

**Data-directed control flow:** inherent part of many complex algorithms, difficult to optimize

Generic implementation strategies:

- **wait** until condition resolved, then start branch execution
- **if-conversion:** pre-compute all possible outcomes, then select actual ones  
*converts control dependencies into data dependencies*
- **restructure** execution for extraction of parallelism
- **predict** most likely condition(s), **flush** computations if mispredicted

```
x = a + b;
if (x > 100) {
    out = x - 10;
} else {
    out = x * x;
}
return out;
```



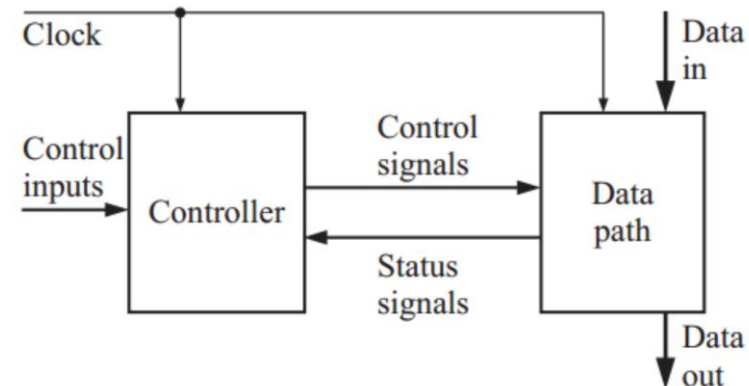

Pure hardware

## Waiting until condition resolved: finite state machines (FSM)


```

x = a + b;
if (x > 100) {
    out = x - 10;
} else {
    out = x * x;
}
return out;

```



## “If-conversion”: pre-computed data multiplexing

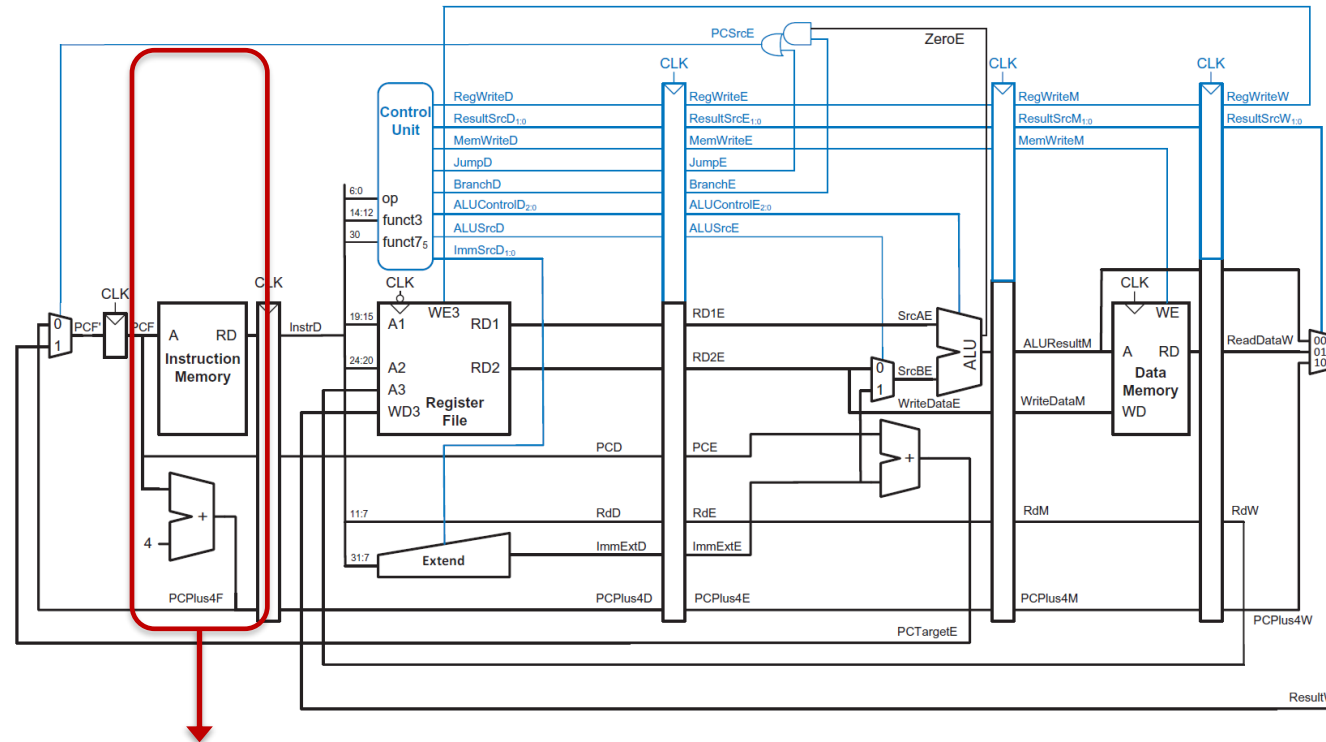


**Bad: redundant computations (unused in active branches)**

## Processor microarchitecture

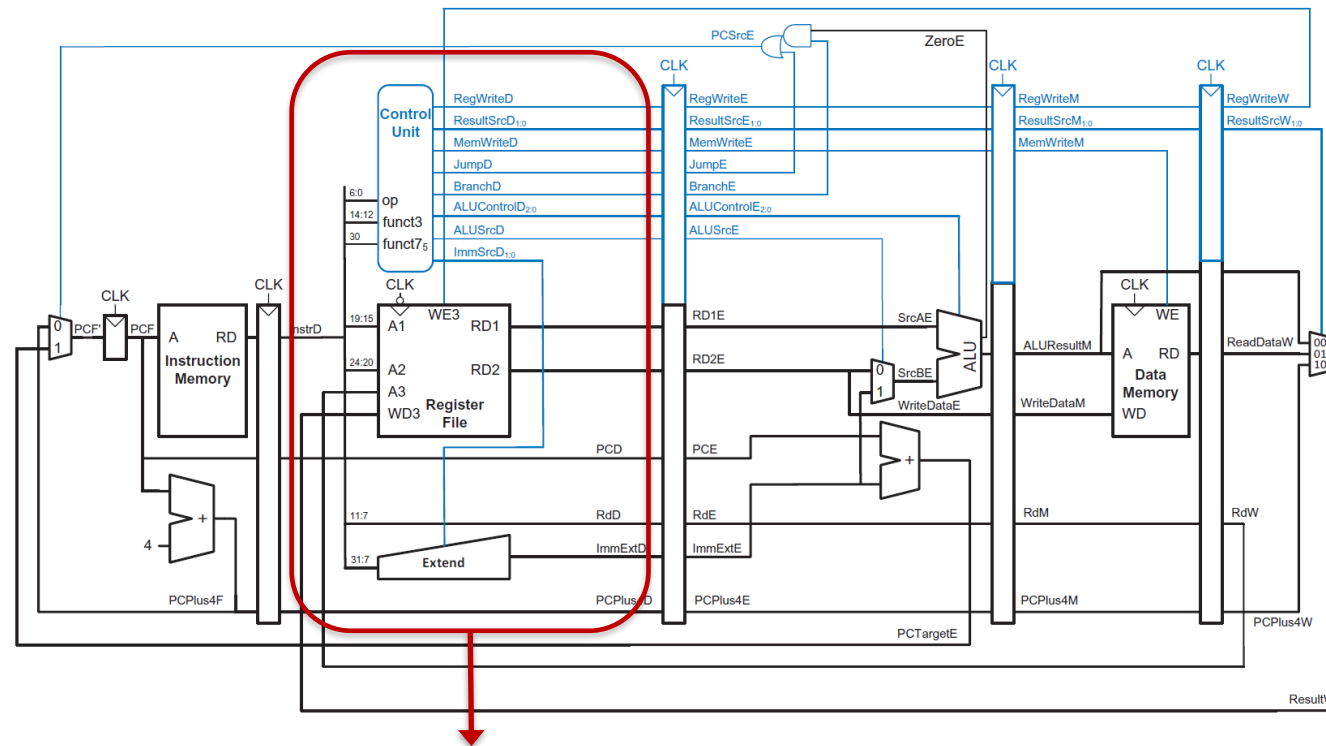


## Branches in classic RISC pipeline (IF stage)



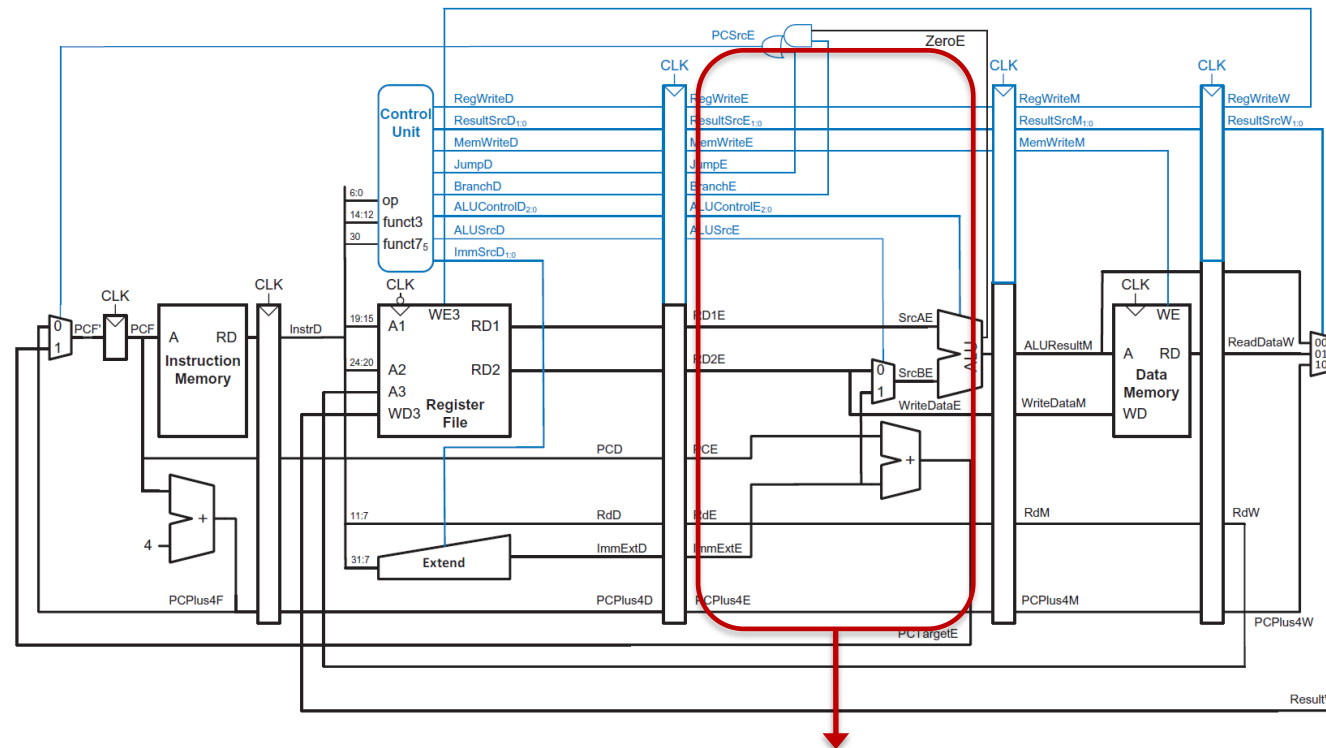
Instruction type (branch?): **unknown**  
 Branch condition (for conditional): **unknown**  
 Static branch target: **unknown**  
 Dynamic (computed) branch target: **unknown**

## Branches in classic RISC pipeline (ID stage)



Instruction type (branch?): **known**  
 Branch condition (for conditional): **unknown**  
 Static branch target: **known**  
 Dynamic (computed) branch target: **unknown**

## Branches in classic RISC pipeline (EX stage)



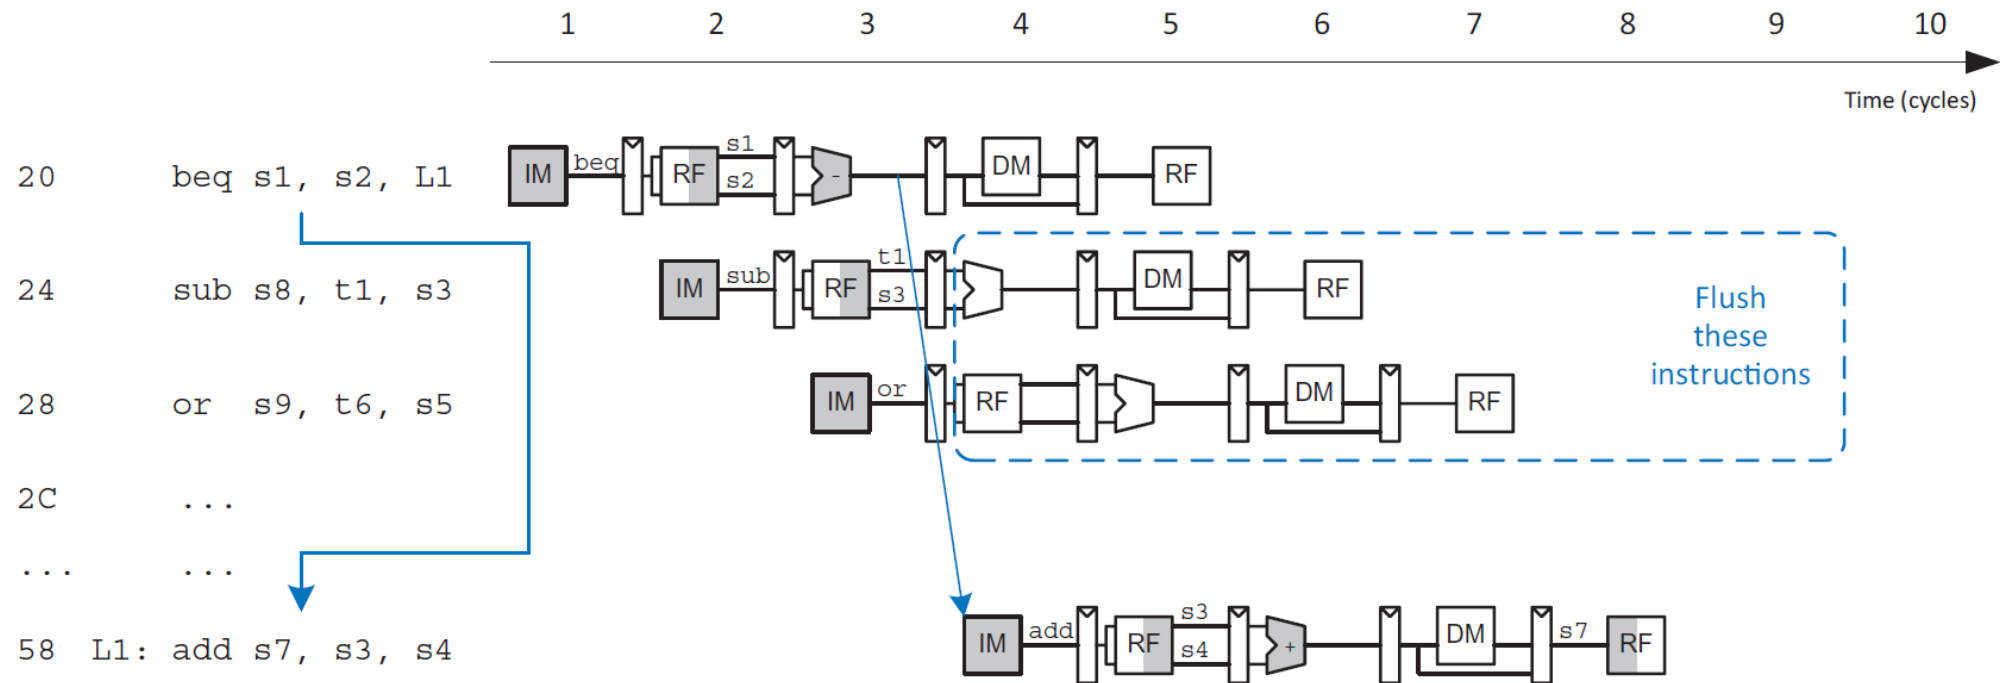
Instruction type (branch?): **known**

Branch condition (for conditional): **known**

Static branch target: **known**

Dynamic (computed) branch target: **known**

## Flushing wasted instructions in classic RISC pipeline



## Instruction flow speculation: dynamic branch prediction


CPU pipelines require **early decision making about instruction control transfers**

dependency: PC write → instruction memory address of the next instruction

But: within instruction flow, there are **uncertainties** on early stages:

- instruction type (branch)?
- dynamic branch conditions?
- computed branch targets?

**Solution: branch prediction**  
(without full awareness of actual program state)



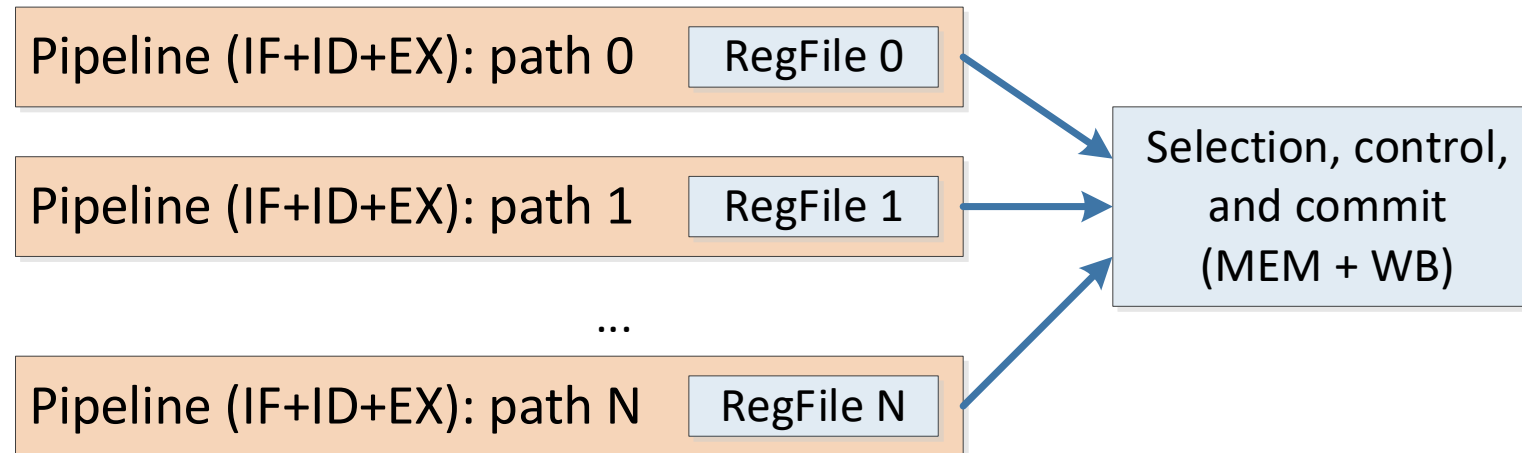
**Branches cannot be predicted with 100% accuracy (e.g. during training, due to random data, etc.),**

**but appear to correlate well with previous branch statistics**

**More information in Lecture 12**

## Multipath execution

Principle: speculatively fetch and execute *several paths simultaneously*, continuously discard false branches



***Bad: redundant computations and storage***

***Issue: nested branching in speculated paths***

## Exceptions and interrupts in processors

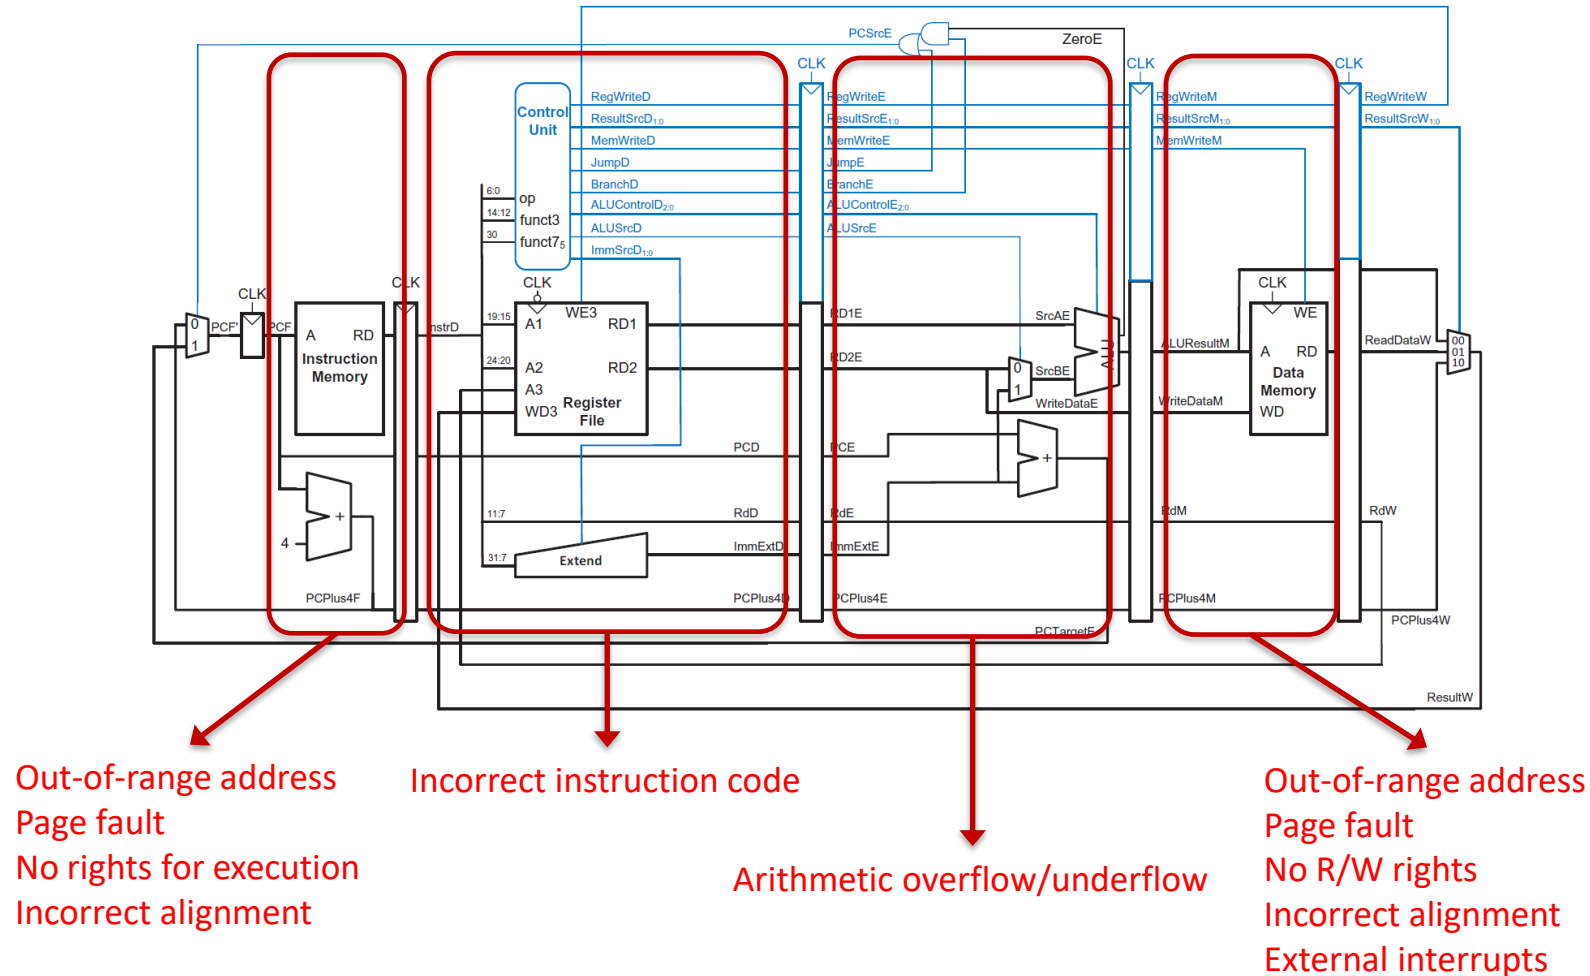
In addition to branches, the following events can alter control flow:

- ***Exception:*** event caused by abnormal program execution
- ***Interrupt:*** event fired by external device

When occurred, control flow is transferred to predefined location where ***software handler*** is expected to reside.

## Exceptions and interrupts in classic RISC pipeline

*Earlier* exceptions have *higher* priority for each instruction





## Processor ISA

## Predicated instructions

Instructions are annotated with ***predicates*** (conditions), instruction with failed predicate = no  
Concept similar to “if-conversion”

Source code	Conventional generated code (with branches)
<pre> if (emp_status == ACTIVE) {     active_emps++;     total_payroll += emp_pay; } else {     inactive_emps++; } </pre>	<pre> {     cmp.ne p1 = rs, ACTIVE // compare emp_status     (p1) br else           // jump to else code if condition fails } .label then {     add rt = rt, rp        // sum total_payroll + emp_pay     add ra = ra, 1         // increment active_emps     br join } .label else {     add ri = ri, 1         // increment inactive_emps } .label join </pre>
<pre> {     cmp.eq p1, p2 = rs, ACTIVE // compare emp_status } {     (p1) add rt = rt, rp        // sum total_payroll + emp_pay     (p1) add ra = ra, 1         // increment active_emps     (p2) add ri = ri, 1         // increment inactive_emps } </pre>	<p>Code with predicated execution</p>

## Restructured branches in ISA: delay slots

ISA feature making ***N instructions after branch execute anyway***

Reduces mispredict penalty

Considered obsolete (difficult to use, difficult to implement in OoO processors)

```
lw r2, r1, 0x100
add r4, r2, r0
beq r1, r10, 0x14
subi r6, r5, 0x20
mul r10, r10, r11
```

will execute anyway

will be delayed after subi

### GCC:

-fdelayed-branch

If supported for the target machine, attempt to reorder instructions to exploit instruction slots available after delayed branch instructions.

Enabled at levels -O1, -O2, -O3, -Os, but not at -Og.

## Prediction hints

Some ISAs allocate **bits** in instructions to **hint** the processor with (un)likely branch occurrence

Reduce misprediction ratio (or can be ignored by processor)

Not widely used in practice

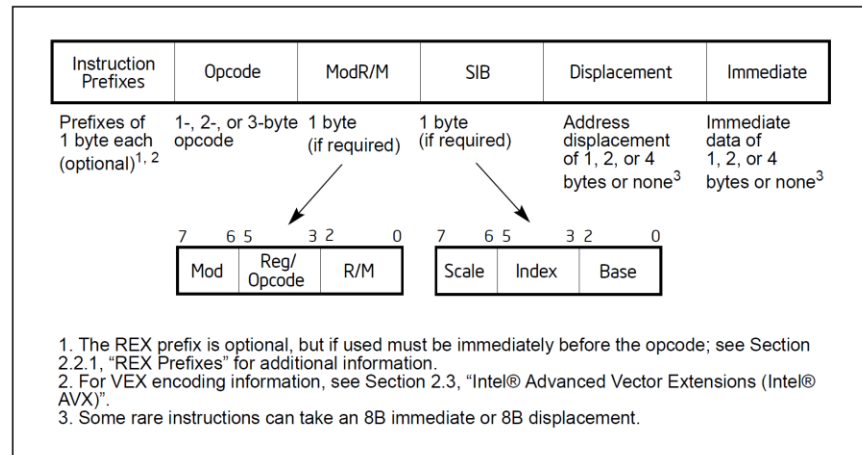


Figure 2-1. Intel 64 and IA-32 Architectures Instruction Format

### INSTRUCTION FORMAT

- BNDCFGU.EN and/or IA32\_BNDCFGS.EN is set.
- When the F2 prefix precedes a near CALL, a near RET, a near JMP, or a near Jcc instruction (see Chapter 17, "Intel® MPX," of the *Intel® 64 and IA-32 Architectures Software Developer's Manual, Volume 1*).
- Group 2
  - Segment override prefixes:
    - 2EH—CS segment override (use with any branch instruction is reserved).
    - 36H—SS segment override prefix (use with any branch instruction is reserved).
    - 3EH—DS segment override prefix (use with any branch instruction is reserved).
    - 26H—ES segment override prefix (use with any branch instruction is reserved).
    - 64H—FS segment override prefix (use with any branch instruction is reserved).
    - 65H—GS segment override prefix (use with any branch instruction is reserved).
  - Branch hints<sup>1</sup>:
    - 2EH—Branch not taken (used only with Jcc instructions).
    - 3EH—Branch taken (used only with Jcc instructions).
- Group 3
  - Operand-size override prefix is encoded using 66H (66H is also used as a mandatory prefix for some instructions).
- Group 4
  - 67H—Address-size override prefix.

## Compilers and programming

## If-conversion in SW: branchless programming

Has branches  
Unfriendly to pipelining

```
if (a > b) {  
    return a;  
} else {  
    return b;  
}
```

No branches  
Friendly to pipelining

```
return (a > b) * a + (a <= b) * b;
```

## Advanced branchless programming: function tables

```
#include <stdio.h>
#include <stdlib.h>

void fun0(int a)
{
    printf("Fun0: a: %d\n", a);
}

void fun1(int a)
{
    printf("Fun1: a: %d\n", a);
}

void fun2(int a)
{
    printf("Fun2: a: %d\n", a);
}

void call_fun(int fun_num, int a) {    // function caller
    switch (fun_num) {
        case 0: fun0(a); break;
        case 1: fun1(a); break;
        case 2: fun2(a); break;
    }
}

int main()
{
    // calling functions
    call_fun(0, 10);
    call_fun(1, 11);
    call_fun(2, 12);

    return 0;
}
```

Many branches  
(in call\_fun function)

```
#include <stdio.h>
#include <stdlib.h>

void fun0(int a)
{
    printf("Fun0: a: %d\n", a);
}

void fun1(int a)
{
    printf("Fun1: a: %d\n", a);
}

void fun2(int a)
{
    printf("Fun2: a: %d\n", a);
}

void (*fun_ptr[3])(int);    // table of functions

void call_fun(int fun_num, int a) {    // function caller
    (*fun_ptr[fun_num])(a);
}

int main()
{
    // initializing a table of functions
    fun_ptr[0] = &fun0;
    fun_ptr[1] = &fun1;
    fun_ptr[2] = &fun2;

    // calling functions
    call_fun(0, 10);
    call_fun(1, 11);
    call_fun(2, 12);

    return 0;
}
```

Less branches


Result:

```
Fun0: a: 10
Fun1: a: 11
Fun2: a: 12
```

## Restructuring code for parallel architectures: loop optimizations

### Loop unrolling

```
for (i = 0; i < 4; i++)
{
    r[i] = a[i] + b[i];
}
```



No Unrolling  
1 Adder shared for 4 additions  
Latency = 4 cycles

```
r[0] = a[0] + b[0];
r[1] = a[1] + b[1];
r[2] = a[2] + b[2];
r[3] = a[3] + b[3];
```



Unrolling = 4 (Full)  
4 Adders in parallel  
Latency = 1 cycle

### Loop merging


```
for (i = 0; i < 32; i++)
{
    a[i] = b[i] * c[i];
}
for (i = 0; i < 16; i++)
{
    z[i] = a[i] + x[i];
}
for (i = 0; i < 32; i++)
{
    atmp = b[i] * c[i];
    if (i < 16)
        z[i] = atmp + x[i];
}
```

No Merging  
Loops execute sequentially  
Latency = 48 cycles


Merging Enabled  
Loops execute in parallel  
Latency = 32 cycles

### Loop pipelining

```
for (i = 0; i < 3; i++)
{
    out[i] = (in[i] * coef1) + coef2;
}
```



No Pipelining  
Latency = 12 cycles  
Throughput = 4 cycles



Pipelining Initiation Interval = 1  
Latency = 6 cycles  
Throughput = 1 cycle



## Software prediction hints (GCC)

```
#define likely(x) __builtin_expect(!!(x), 1)
#define unlikely(x) __builtin_expect(!!(x), 0)

if (likely(x > 0)) {
    // This block is expected to execute more frequently
} else {
    // This block is less likely to be executed
}

if (unlikely(error_condition)) {
    // Handle error - expected to be executed infrequently
}
```

***Generated machine code will make more likely branches more likely to predict as taken by CPU***

## Summary of control flow implementations

	Hardware	Processor microarchitecture	Processor ISA	Software
<b>Waiting until condition resolved</b>	Finite state machines	Non-pipelined implementations	Most conventional ISAs	Most conventional software
<b>If-conversion</b>	Data pre-computation and multiplexing	Multipath execution	Predicated instructions	If-conversion in compilers, branchless programming
<b>Restructuring execution</b>	Custom hardware parallelization optimizations	Out-of-order execution	Delay slots	Compiler (e.g. loop) optimizations
<b>Prediction</b>	Prediction in hardware logic	Branch prediction	ISA prediction hints	Software prediction hints



**Thank you for the lesson!**

Alexander Antonov, Assoc. Prof., [antonov@itmo.ru](mailto:antonov@itmo.ru)

Hangzhou, 2025