



**Computer Systems Design**

Lesson 12

Dynamic branch prediction

Open-source superscalar OoO CPU cores

Alexander Antonov, Assoc. Prof., ITMO University

Hangzhou, 2025

## Outline the lesson

- Branch prediction goals
- Branch predictors
  - Branch target buffer
  - Return address stack
  - Adaptive branch predictors
  - TAGE
  - Neural branch predictors
  - Combining branch predictors
- Open-source superscalar OoO CPU cores

## Dynamic branch prediction

## Branch prediction: what and how to predict?

Data to predict:

- 1) branch outcome (taking/not taking branch)
- 2) branch target address

Typically used information:

- current instruction address (PC)
- cached branch target addresses
- branch history (shift register with previous taken/not taken outcomes)
- accumulated statistics for current PC/history combinations

To reduce resource requirements for statistics accumulation, instruction address and history are **hashed**

Problem: **interference** between aliased branches

## Static branch prediction

### Never branch

- effectively not making any predictions
- simplest to implement, but performs badly

### Backwards taken/forwards not-taken (BTFNT)

- “backwards taken” - good for loops (>90% precision)
- other branches – insufficient (50% precision)

***In total, static schemes have insufficient performance (70-80% precision)***

***Adaptation to runtime needed (>99% achievable)***

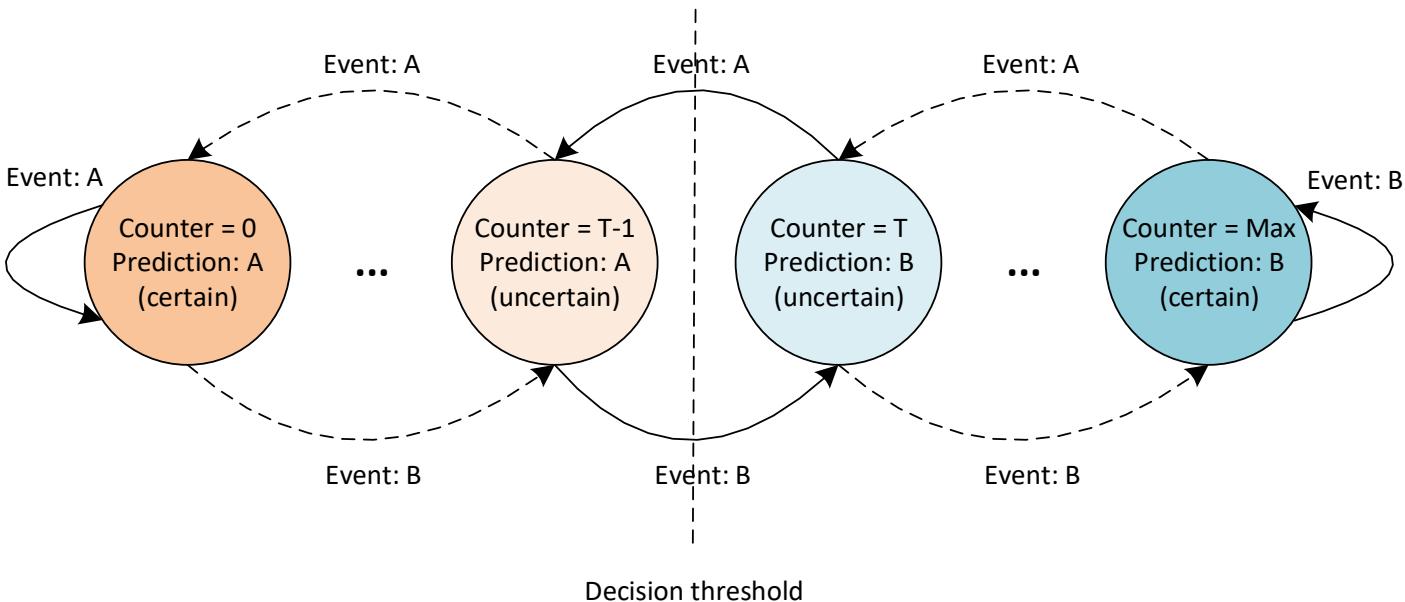
# Statistics accumulation: saturation arithmetic

Principle: **overflow** is replaced with **saturation**

*min/max values preserved*

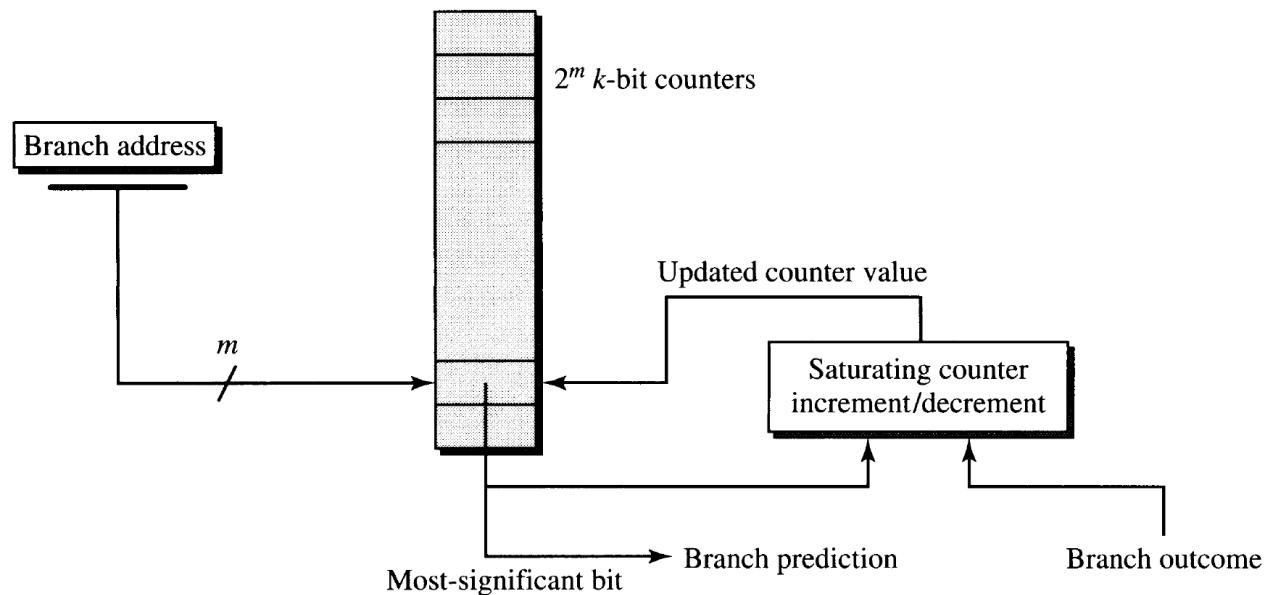
e.g. max value: 0x7F;  $0x70 + 0x20 = 0x7F$  (*not 0x10 or 0x90*)

Prediction is based on accumulated statistics with increasing confidence for monotonic events (hysteresis, inertia)



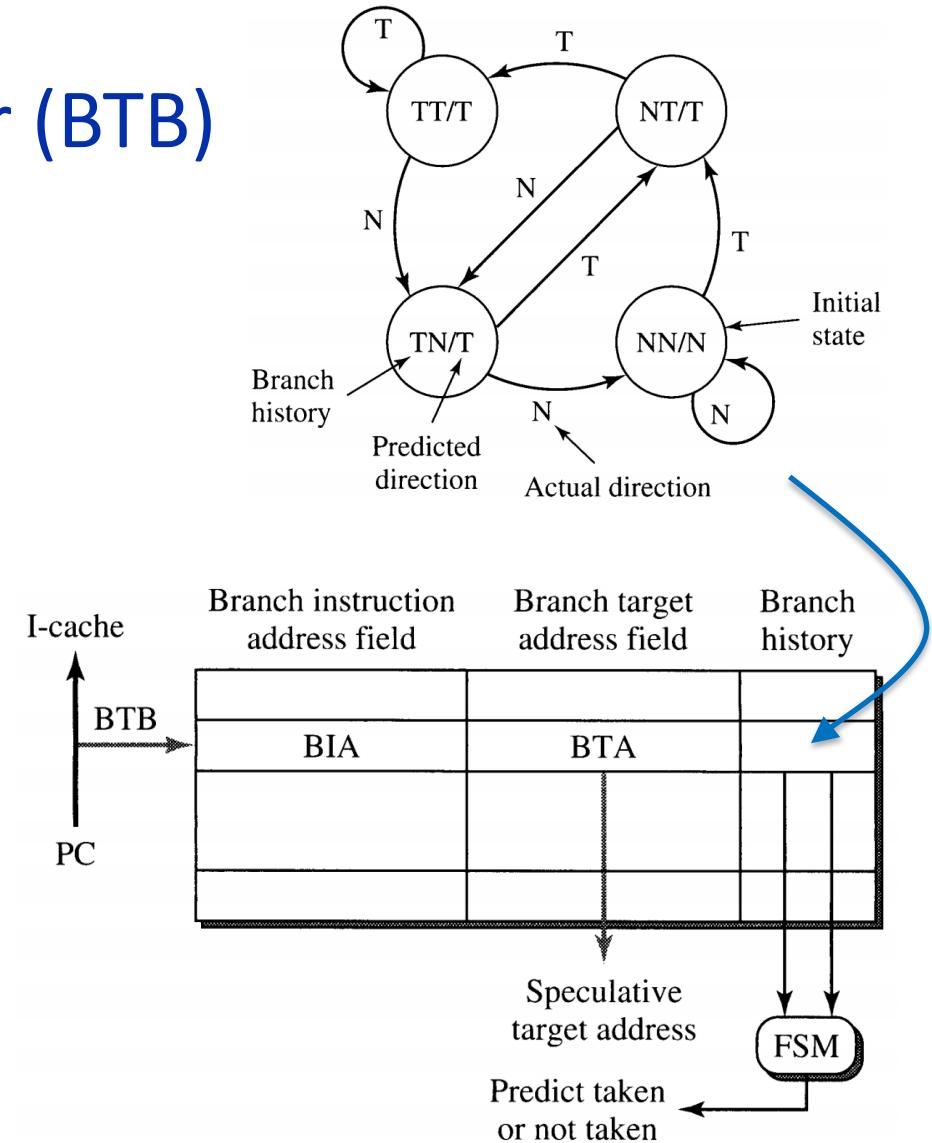
## Basic branch outcome predictor: Smith's predictor (1981)

- branch history is implemented as saturating counters
- accessed by instruction identifier (PC)
- MSB of branch counter is used as prediction

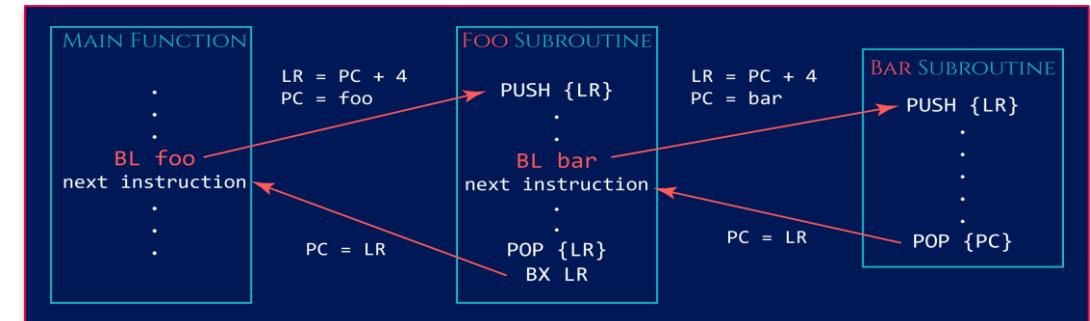
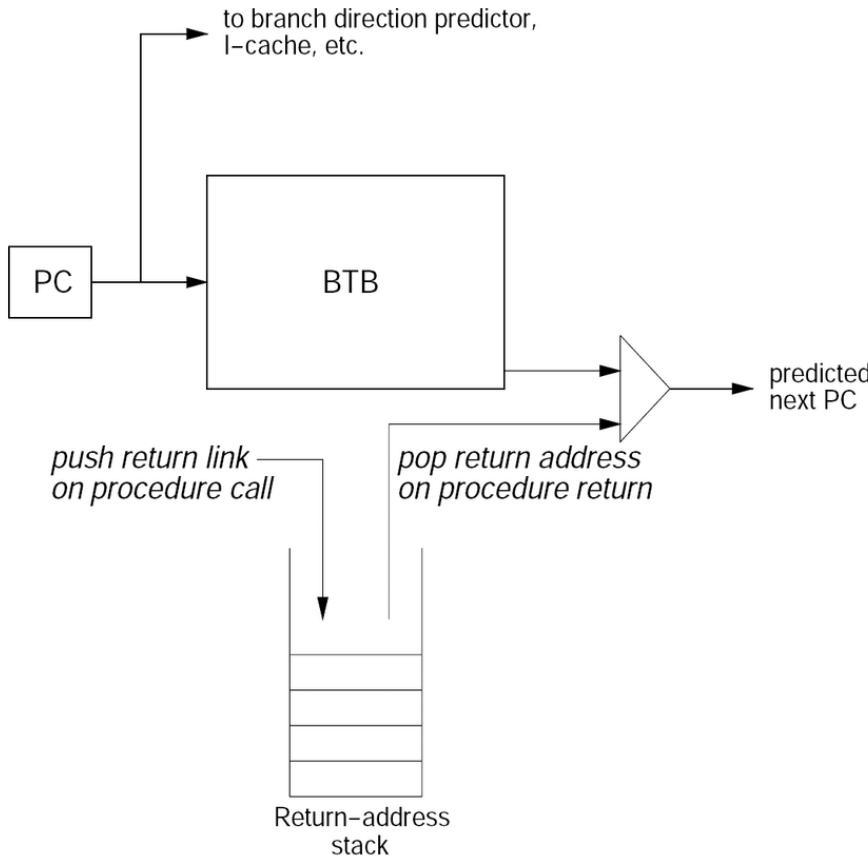


## Target prediction: Branch Target Buffer (BTB)

- cache for taken branches
- branch history is implemented as saturating counter (usually biased towards taking)
- accessed by instruction identifier (PC)
- if branch predicted taken: branch target fetched and used as PC for next instruction



# Return Address Stack (RAS)

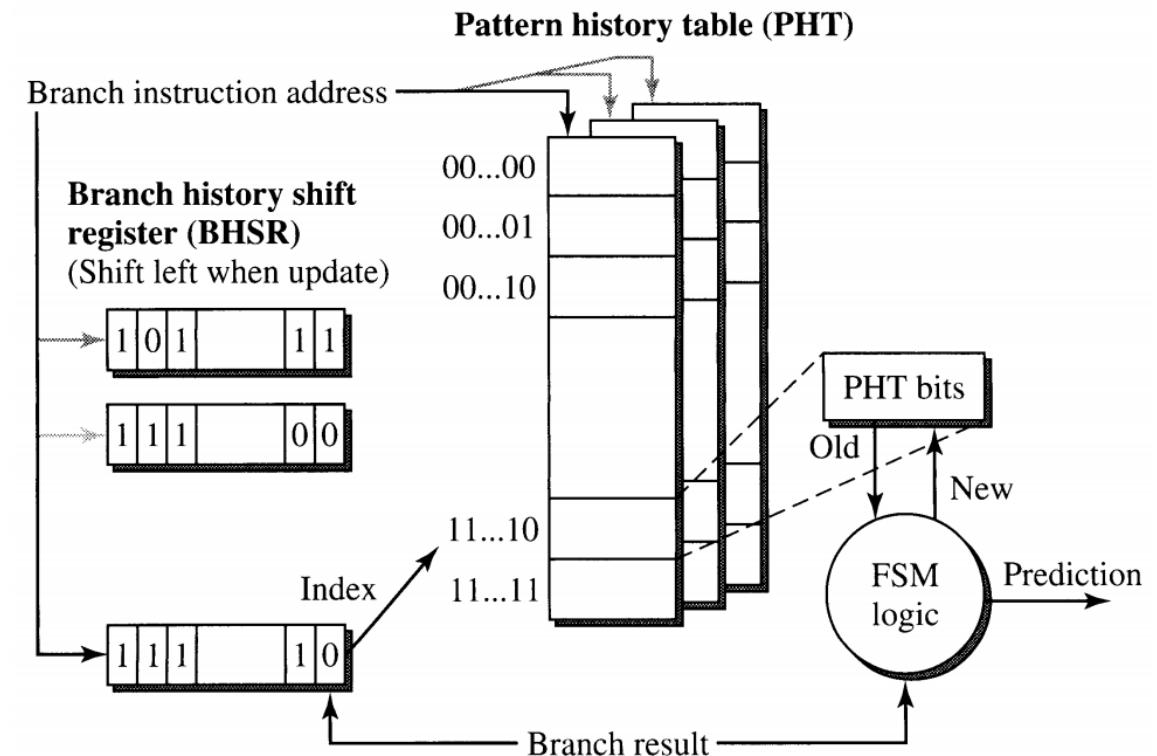


- Function: stack saving return link when calling procedures
- Activates: when returning from procedure
- Prediction: return from procedure will take place to instruction right after the call

## Two-level adaptive branch predictor (1991)

**Observation:** branch behavior correlates well with previous branch patterns

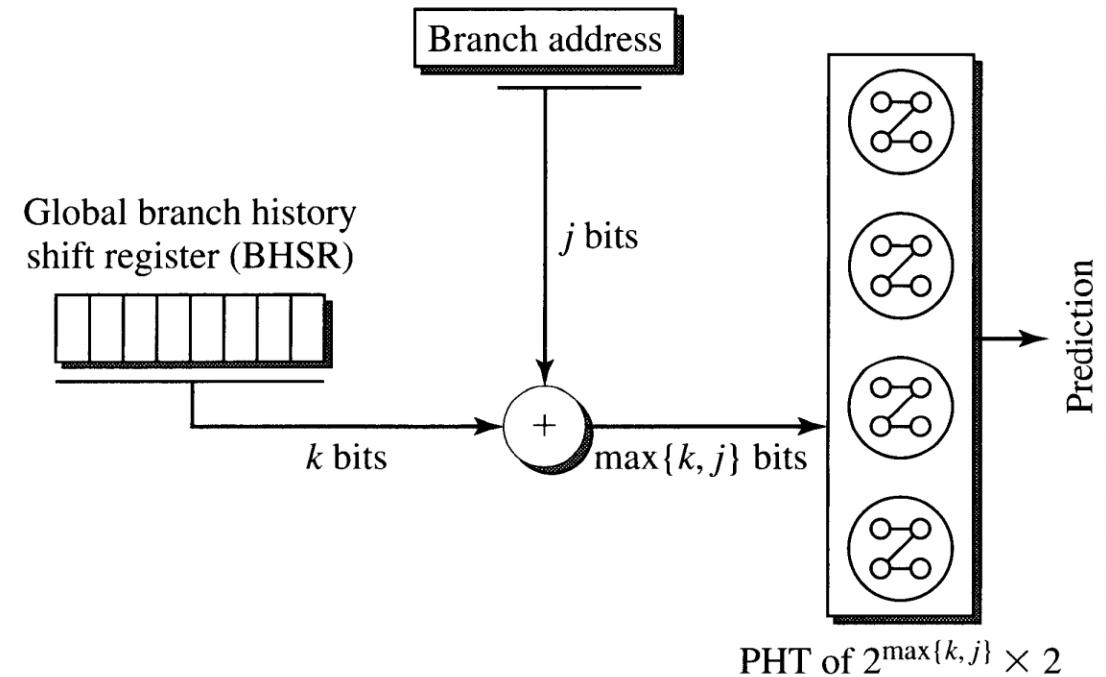
- **Idea:** *selective* statistics accumulation for *branch instruction and branch history*
- PHT is a 2D table (branch instruction x branch history)
- BHSR can be global (one) or multiple individual for each branch (in picture)
- >95% prediction accuracy achievable



## Gshare (McFarling, 1993)

**Observation:** correlated predictor can be *simplified* without significant loss of quality

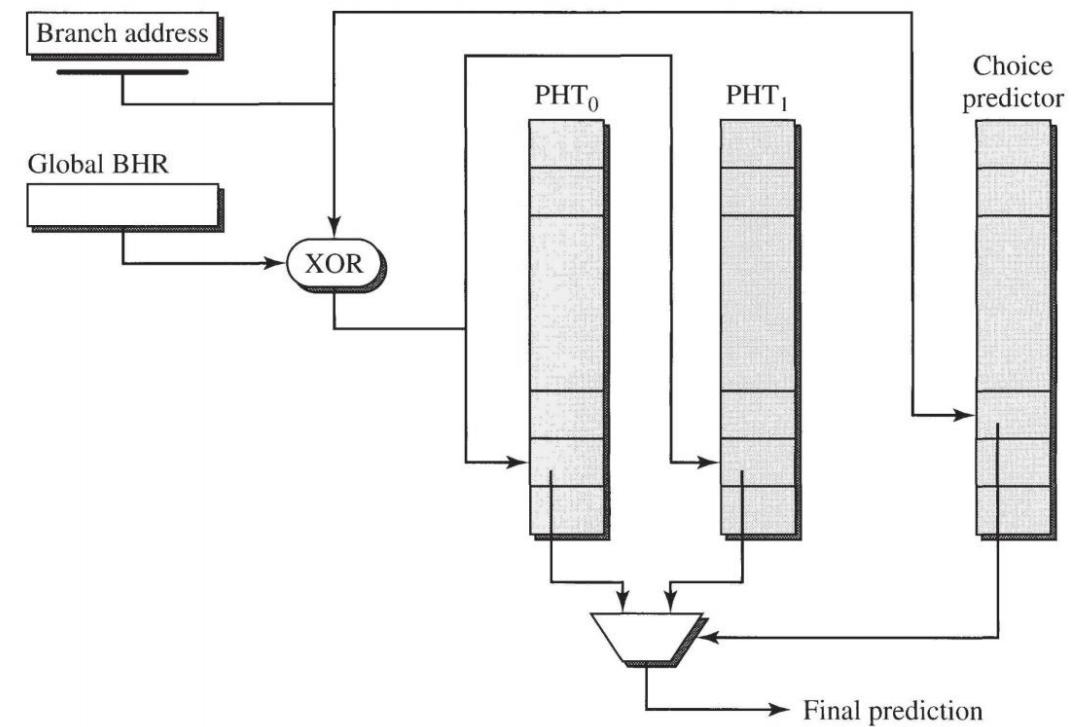
- Idea: global BHSR is *hashed* (XOR) with least significant bits of branch address
- Selectivity of branch address / BHSR is partially preserved  
*issue: collisions (tolerable)*
- Offers good complexity/ accuracy trade-off



## Bi-mode predictor (1997)

**Observation:** some branches have strong “taken” bias, others – “non-taken”.  
For these different branches aliases are destructive.

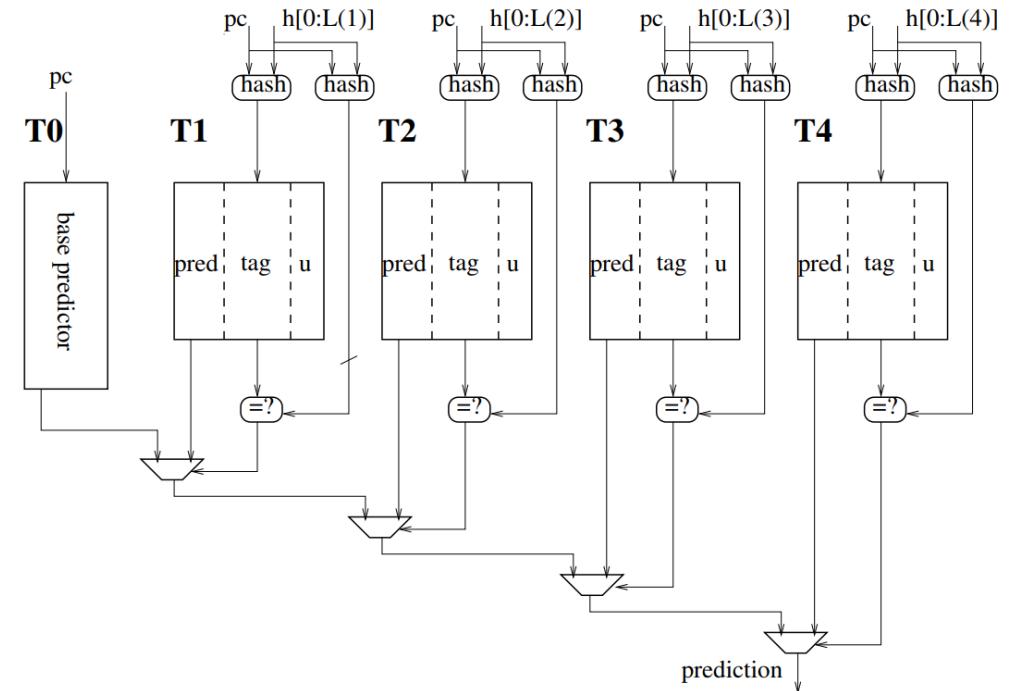
- Branches with “taken” and “non-taken” bias are placed to different PHTs
- Choice predictor (saturating counter) selects what PHT to use
- Destructive interference is reduced



## TAGE (TAgged GEometric history length, 2006)

**Observation:** *high-confidence* predictions based on the *longest history* are the most reliable

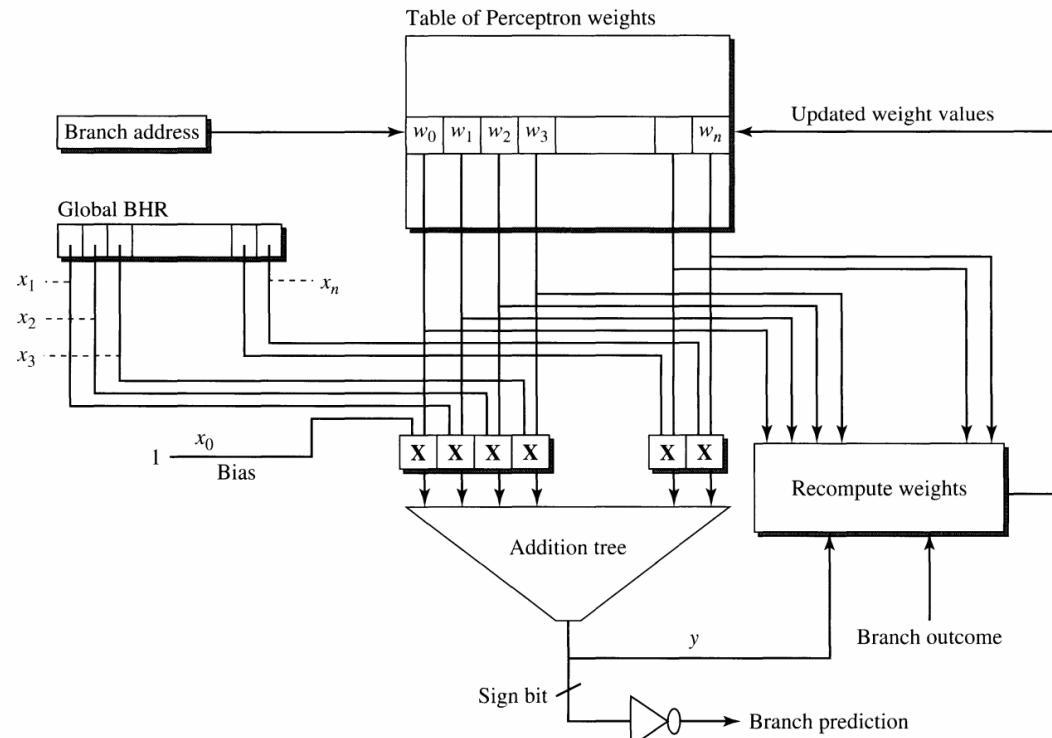
- Series of predictors accumulating predictions for *geometrically increasing history length*
- *High-confidence* prediction *based on the longest history* is used
- Considered one of the best available solutions



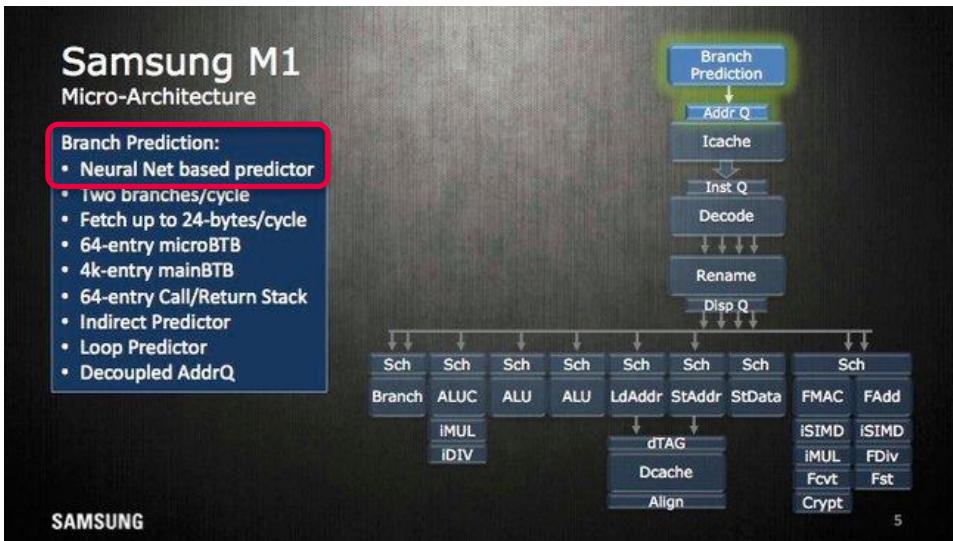
## Neural predictor (2000s ...)

**Observation:** various bits in branch history register have ***various importance*** (“***weights***”) in decision making

- Coefficient vector defines weights of bits in BHR
- Perceptron makes prediction based on weight vector and continuously updates it (“***learns***”) based on actual branch outcomes

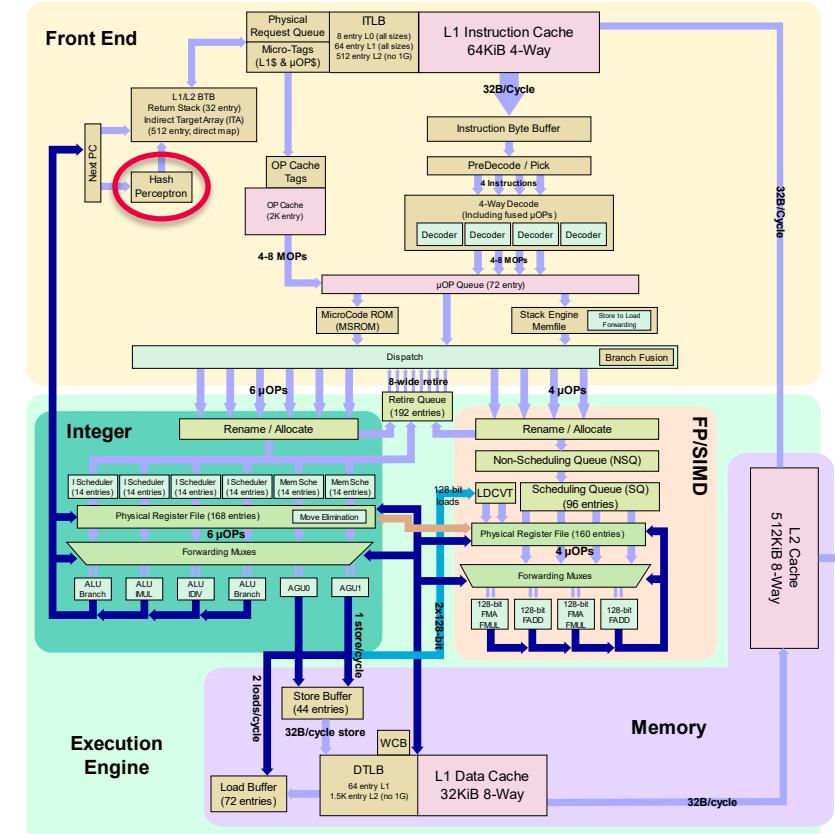


# Neural predictors in modern CPU cores



*Samsung Exynos*

[https://www.theregister.com/2016/08/22/samsung\\_m1\\_core/](https://www.theregister.com/2016/08/22/samsung_m1_core/)

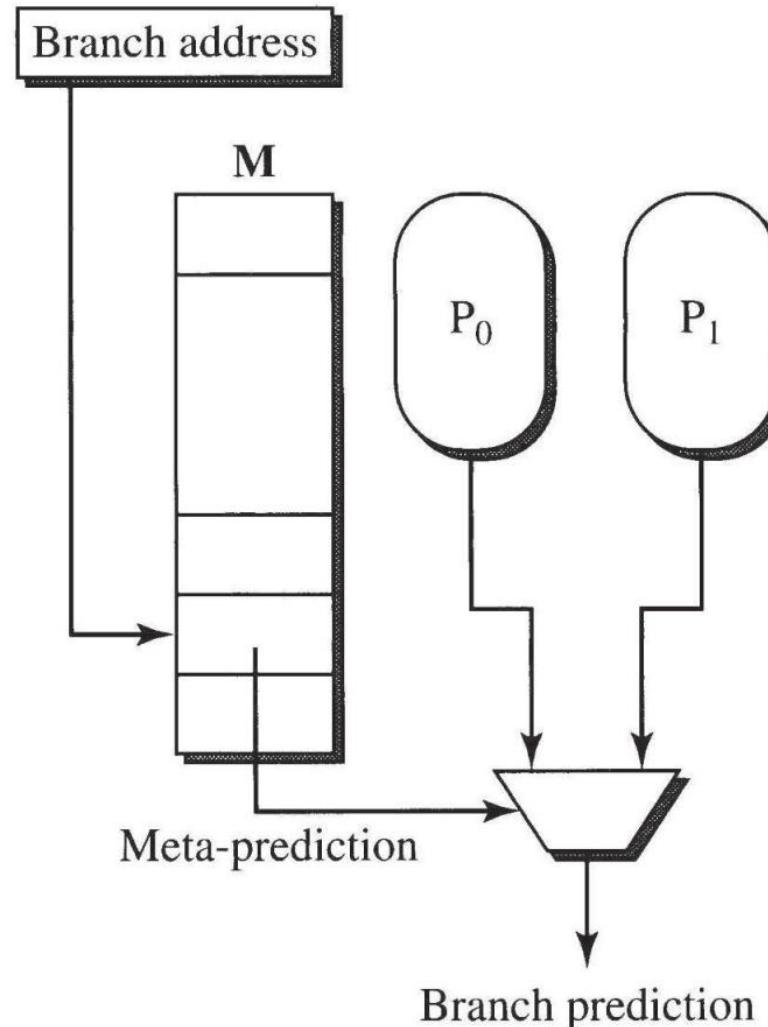


*AMD Zen*

<https://en.wikichip.org/wiki/amd/microarchitectures/zen>

## Hybrid predictors

- several different predictors work in parallel
- effective predictor in current state is chosen based on *meta-predictor*



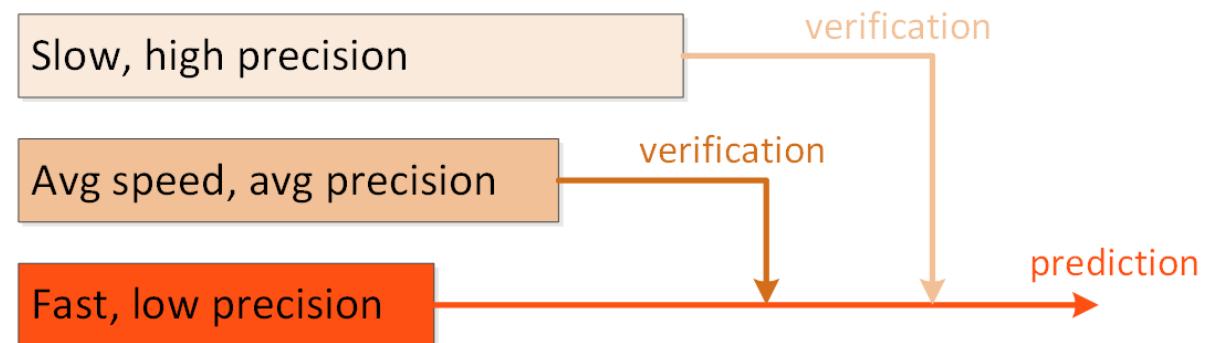
Scott McFarling. Combining Branch Predictors. TN-36, Digital Equipment Corporation, Western Research Laboratory, June 1993

## Hierarchical predictors

**Observation:** precise predictors typically have longer latencies

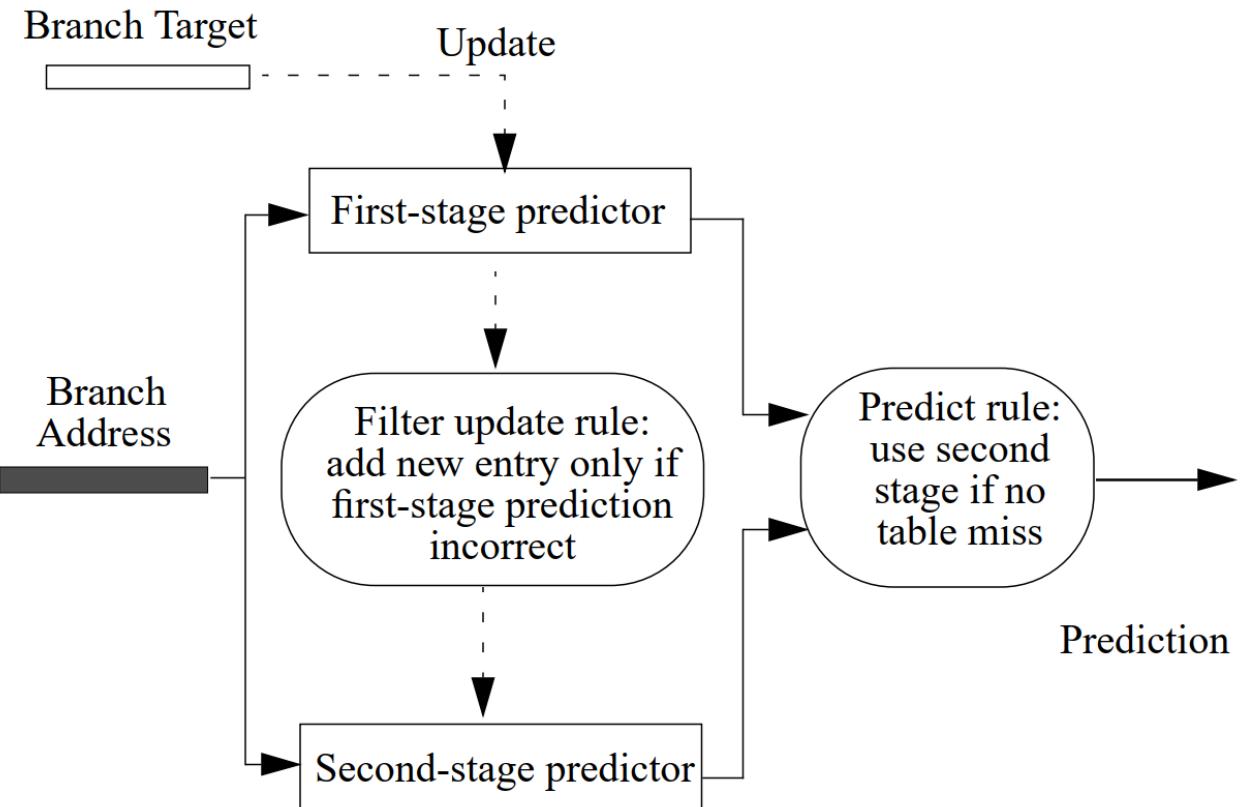
Solution: **combine** fast (but imprecise) and precise (but slow) predictors

- Imprecise predictions output ***first***, precise ones – ***afterwards***
- If precise predictors contradict imprecise ones – ***flush***  
*additional penalty is low (latency margin)*



## Cascaded predictors

- Combines simple and complex predictors
- Classifies branches according to their performance on simple predictor
- If simple predictor performs poorly – entry in complex predictor allocated
- ***Easily predicted branches are filtered – don't overload complex predictor***

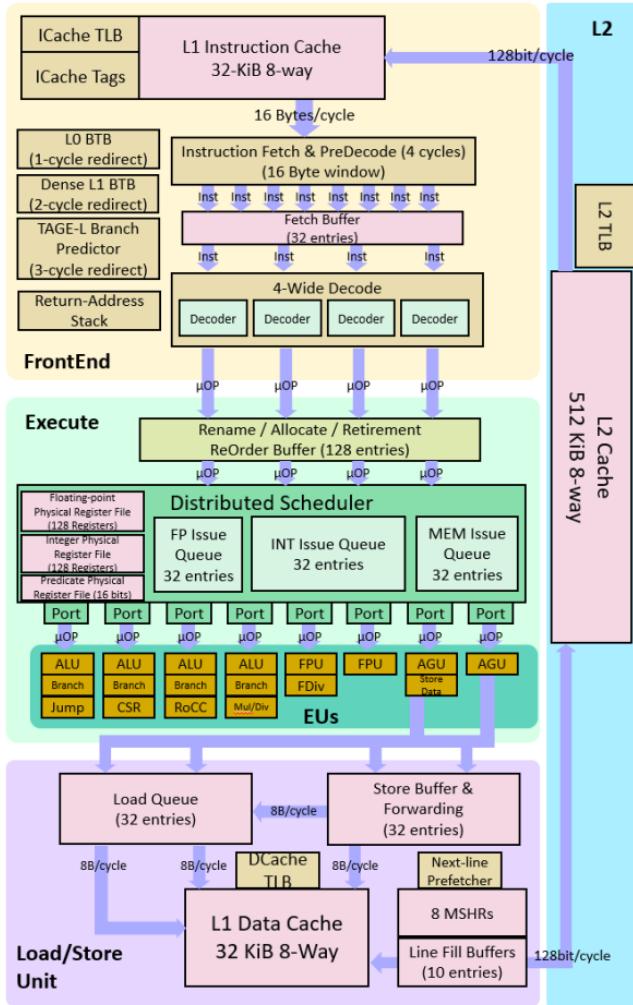


K. Driesen and U. Hözle, "The Cascaded Predictor: Economical and Adaptive Branch Target Prediction", MICRO-31, pp. 249-258, 1998.

Open-source superscalar OoO CPU cores

SonicBOOM/Ocelot

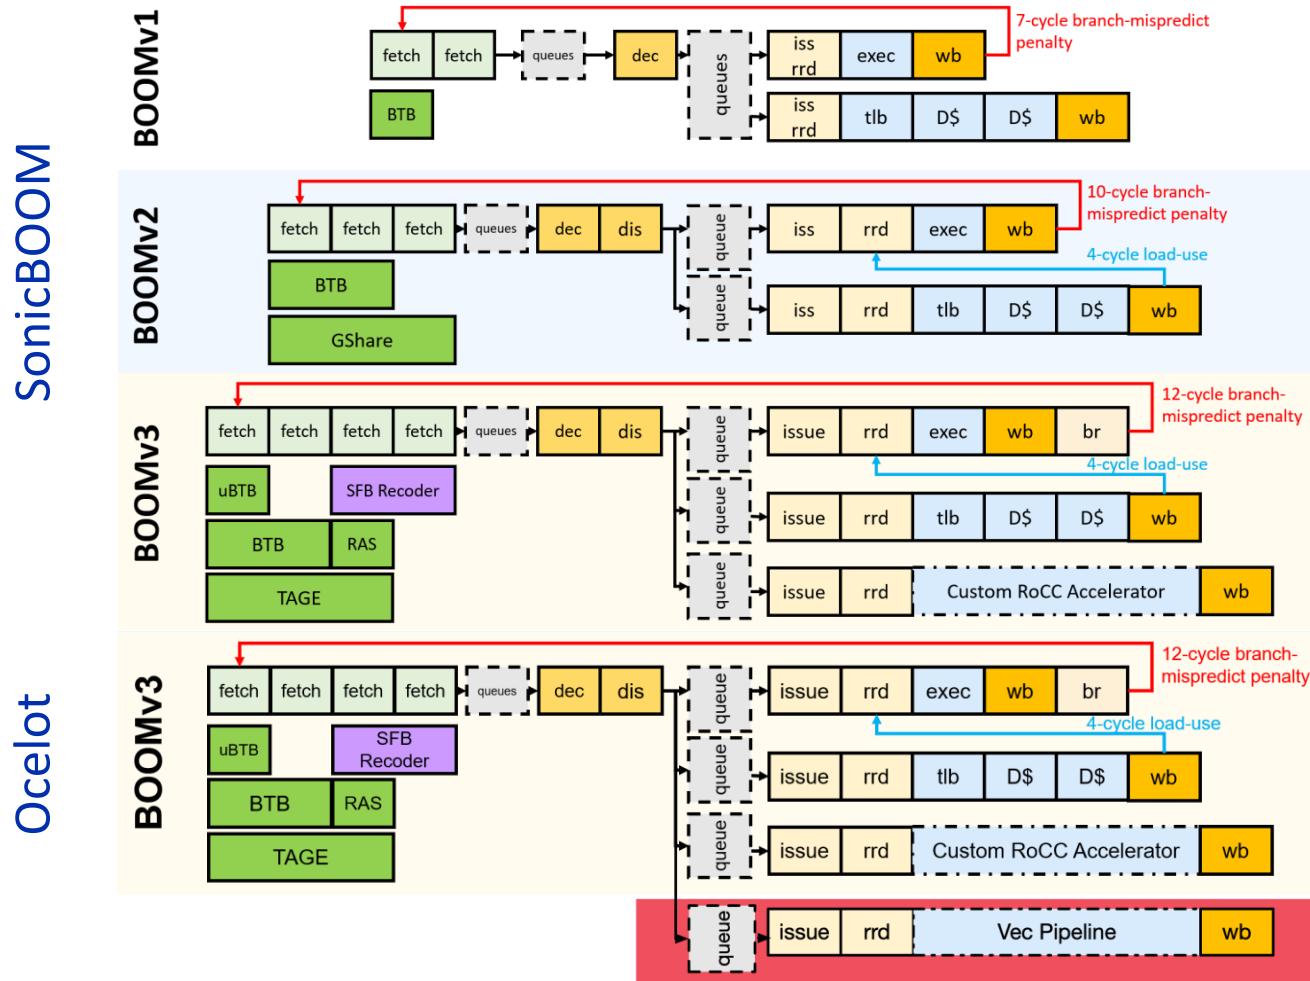
# SonicBOOM/Ocelot (Berkeley/Tenstorrent): design overview



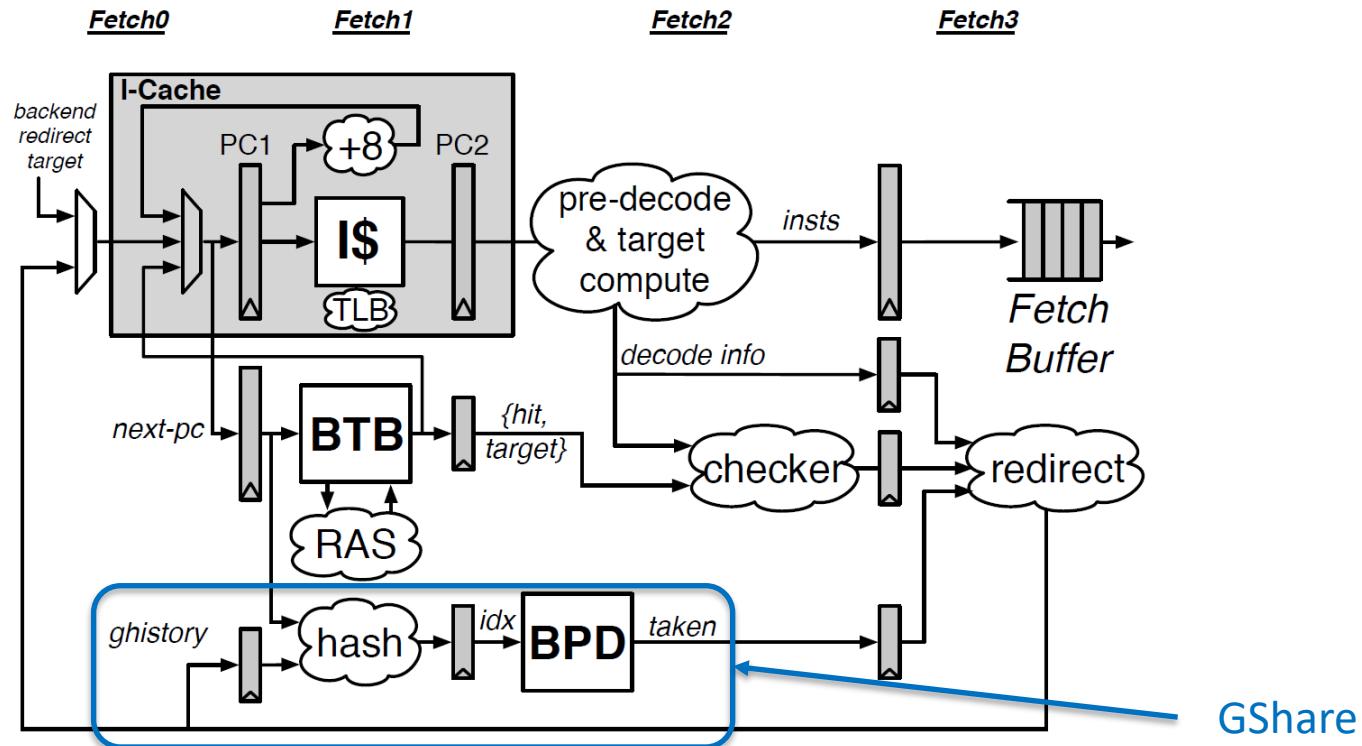
- Three main iterations implemented
- RISC-V architecture (RV64GC, Ocelot: RV64GCV)
- Boots Linux
- Written in Chisel high-level HDL
- Declared as the fastest open source core (x2, 2020), 6.2 CoreMark/MHz
- Demonstrates actual microarchitectural mechanisms *superscalar, speculations, branch prediction, caching, etc.*

Jerry Zhao et al. SonicBOOM: The 3rd Generation Berkeley Out-of-Order Machine. CARRV 2020  
 Repo: <https://github.com/riscv-boom/riscv-boom>, <https://github.com/tenstorrent/riscv-ocelot>  
 Documentation: <https://docs.boom-core.org/en/latest/>

# BOOM v1-3 pipeline structures

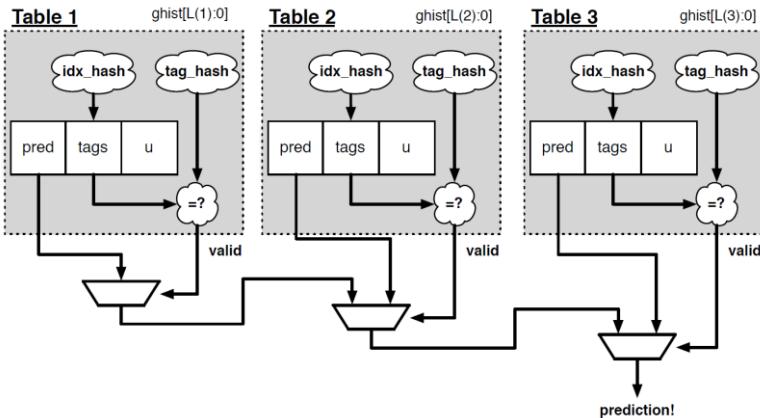


## Instruction fetch: overview (BP: L1 + L2)



Christopher Celio. A Highly Productive Implementation of an Out-of-Order Processor Generator. Technical Report No. UCB/EECS-2018-151. Berkeley, CA. December 2018.

# L3 branch predictor: TAGE

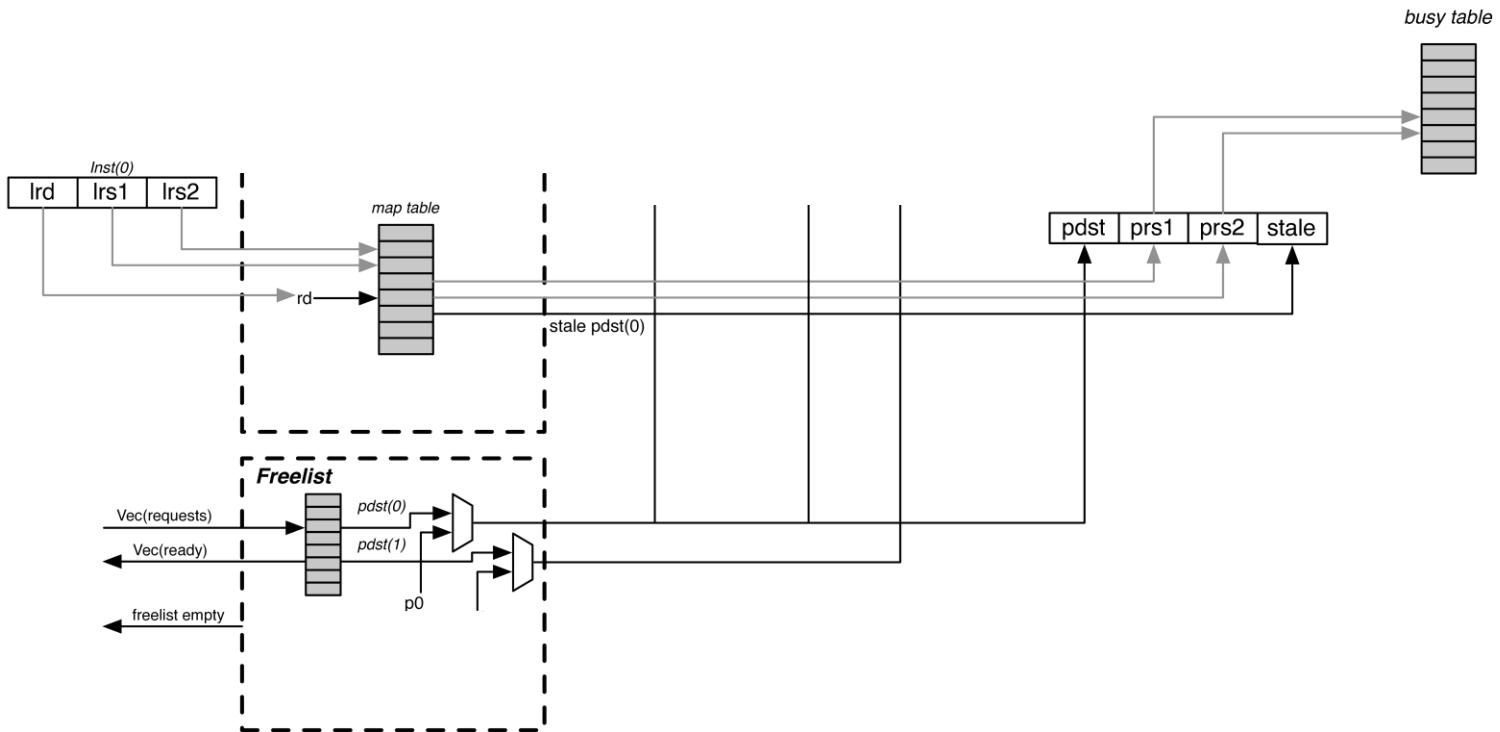


```

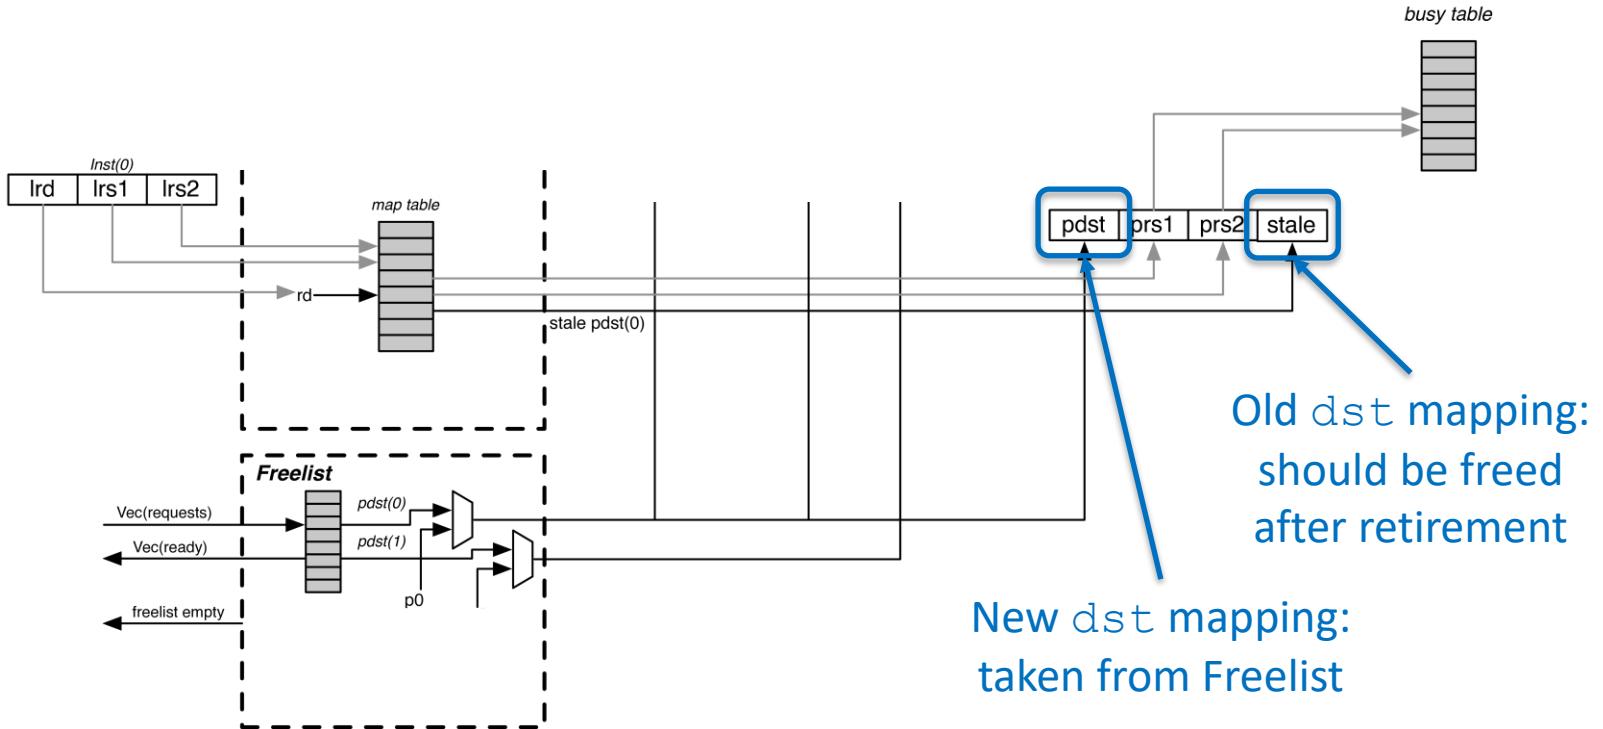
23
24 class TageTable(val nRows: Int, val tagSz: Int, val histLength: Int, val uBitPeriod: Int)
25   (implicit p: Parameters) extends BoomModule()(p)
26   with HasBoomFrontendParameters
27 {
28   require(histLength <= globalHistoryLength)
29
30   val nWrBypassEntries = 2
31   val io = IO(new Bundle {
32     val f1_req_valid = Input(Bool())
33     val f1_req_pc  = Input(UInt(vaddrBitsExtended.W))
34     val f1_req_ghist = Input(UInt(globalHistoryLength.W))
35
36     val f3_resp = Output(Vec(bankWidth, Valid(new TageResp())))
37
38     val update_mask    = Input(Vec(bankWidth, Bool()))
39     val update_taken   = Input(Vec(bankWidth, Bool()))
40     val update_alloc   = Input(Vec(bankWidth, Bool()))
41     val update_old_ctr = Input(Vec(bankWidth, UInt(3.W)))
42
43     val update_pc     = Input(UInt())
44     val update_hist   = Input(UInt())
45
46     val update_u_mask = Input(Vec(bankWidth, Bool()))
47     val update_u      = Input(Vec(bankWidth, UInt(2.W)))
48   })
49
50   def compute_folded_hist(hist: UInt, l: Int) = {
51     val nChunks = (histLength + 1 - 1) / l
52     val hist_chunks = (0 until nChunks) map {i =>
53       hist(min((i+1)*l, histLength)-1, i*l)
54     }
55     hist_chunks.reduce(_.^_)
56   }
57
58   def compute_tag_and_hash(unhashed_idx: UInt, hist: UInt) = {
59     val idx_history = compute_folded_hist(hist, log2Ceil(nRows))
60     val idx = (unhashed_idx ^ idx_history)(log2Ceil(nRows)-1, 0)
61   }
62 }
```

<https://github.com/riscv-boom/riscv-boom/blob/master/src/main/scala/ifu/bpd/tage.scala>

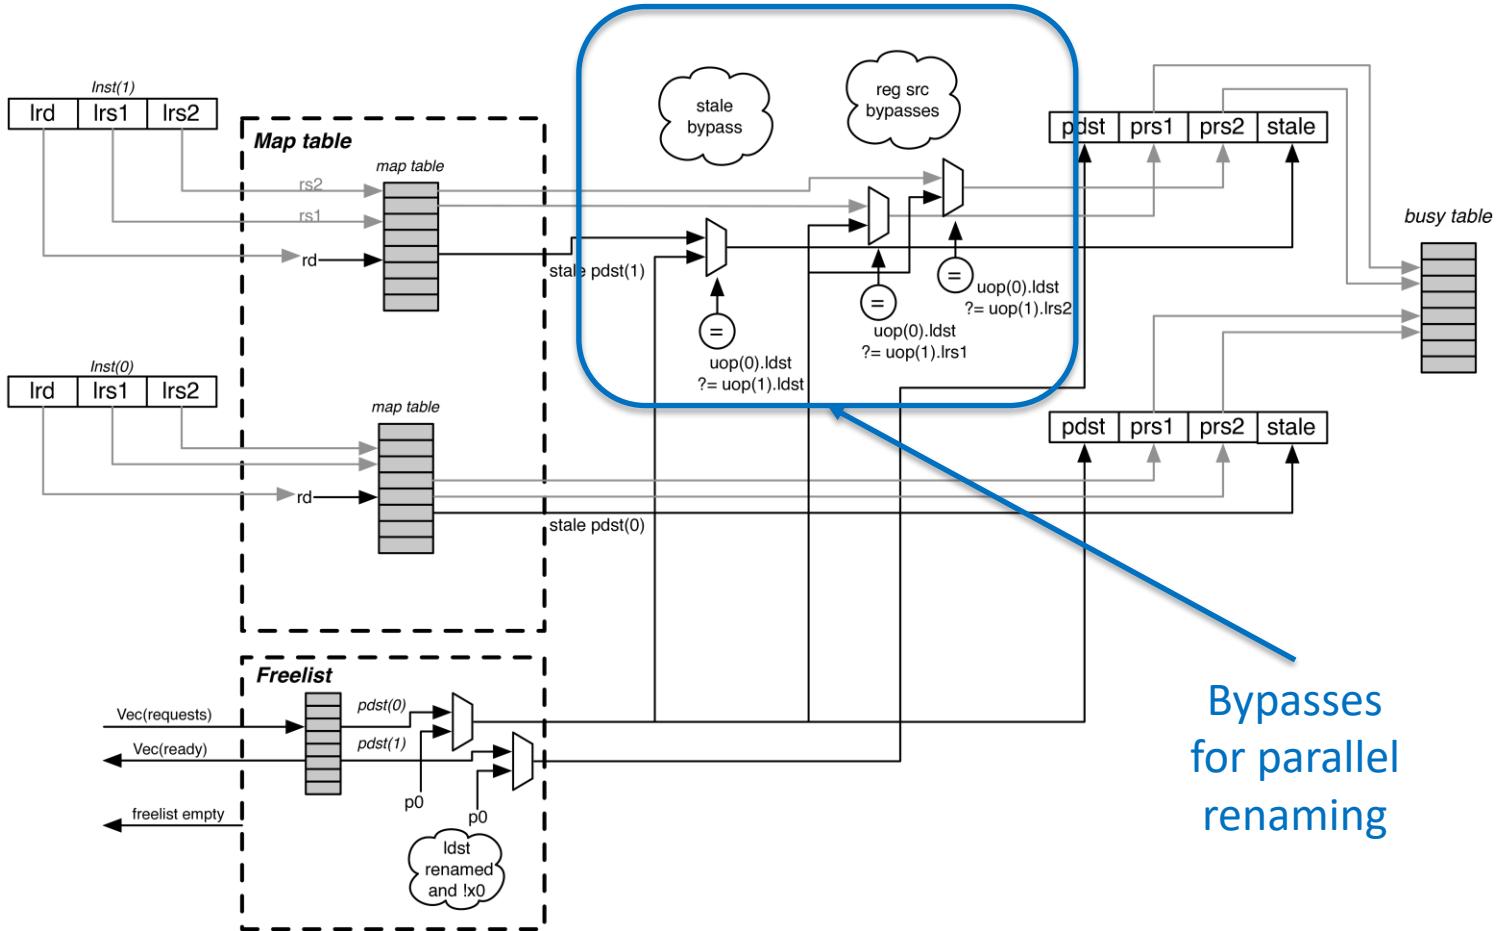
## Register renaming (one channel)



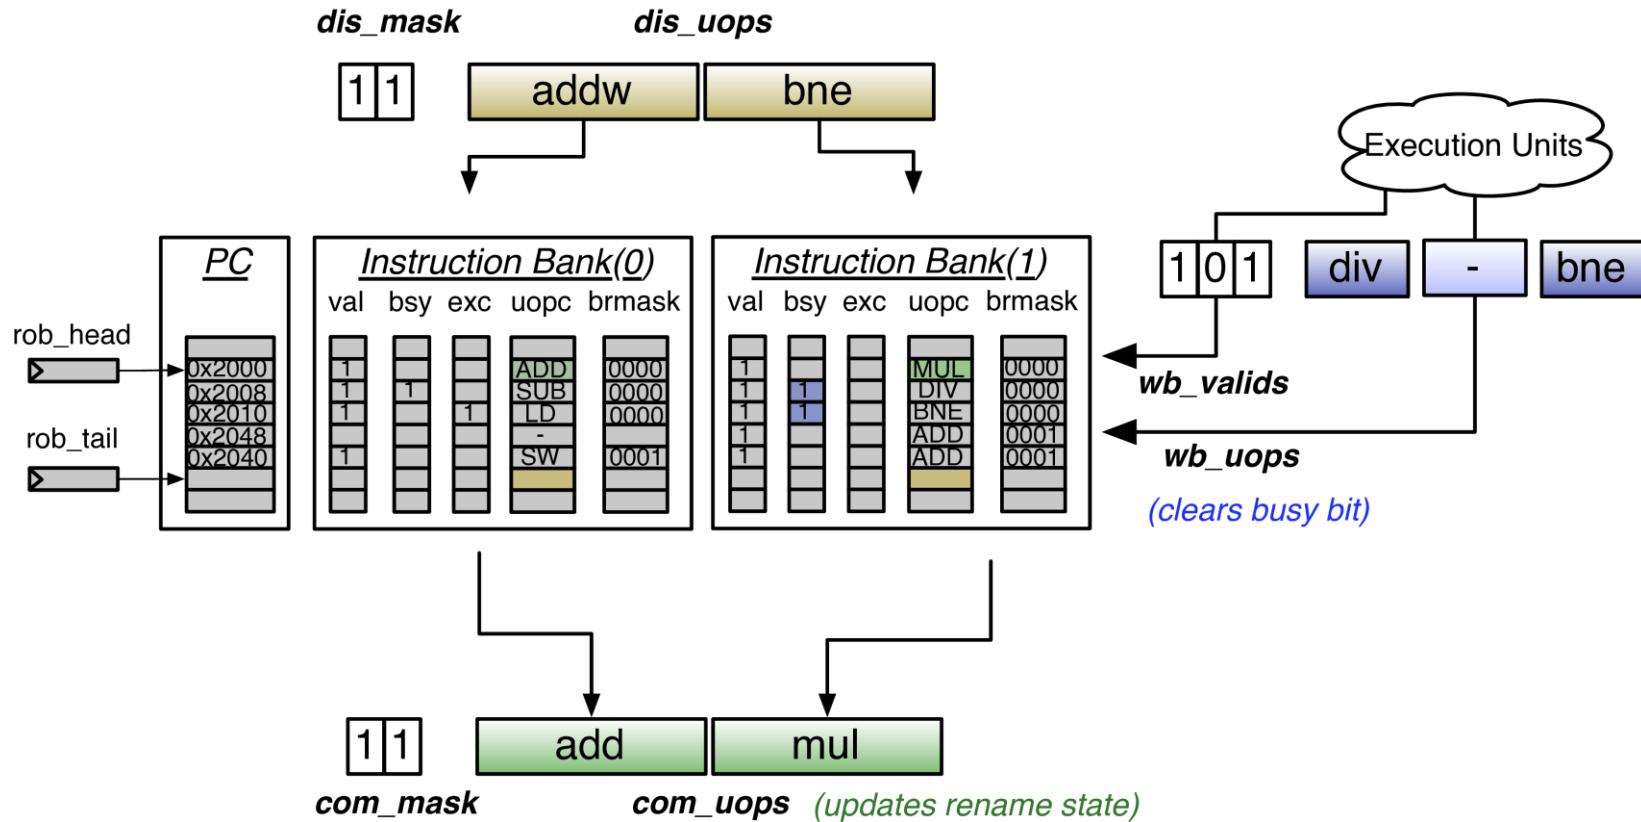
## Register renaming (one channel)



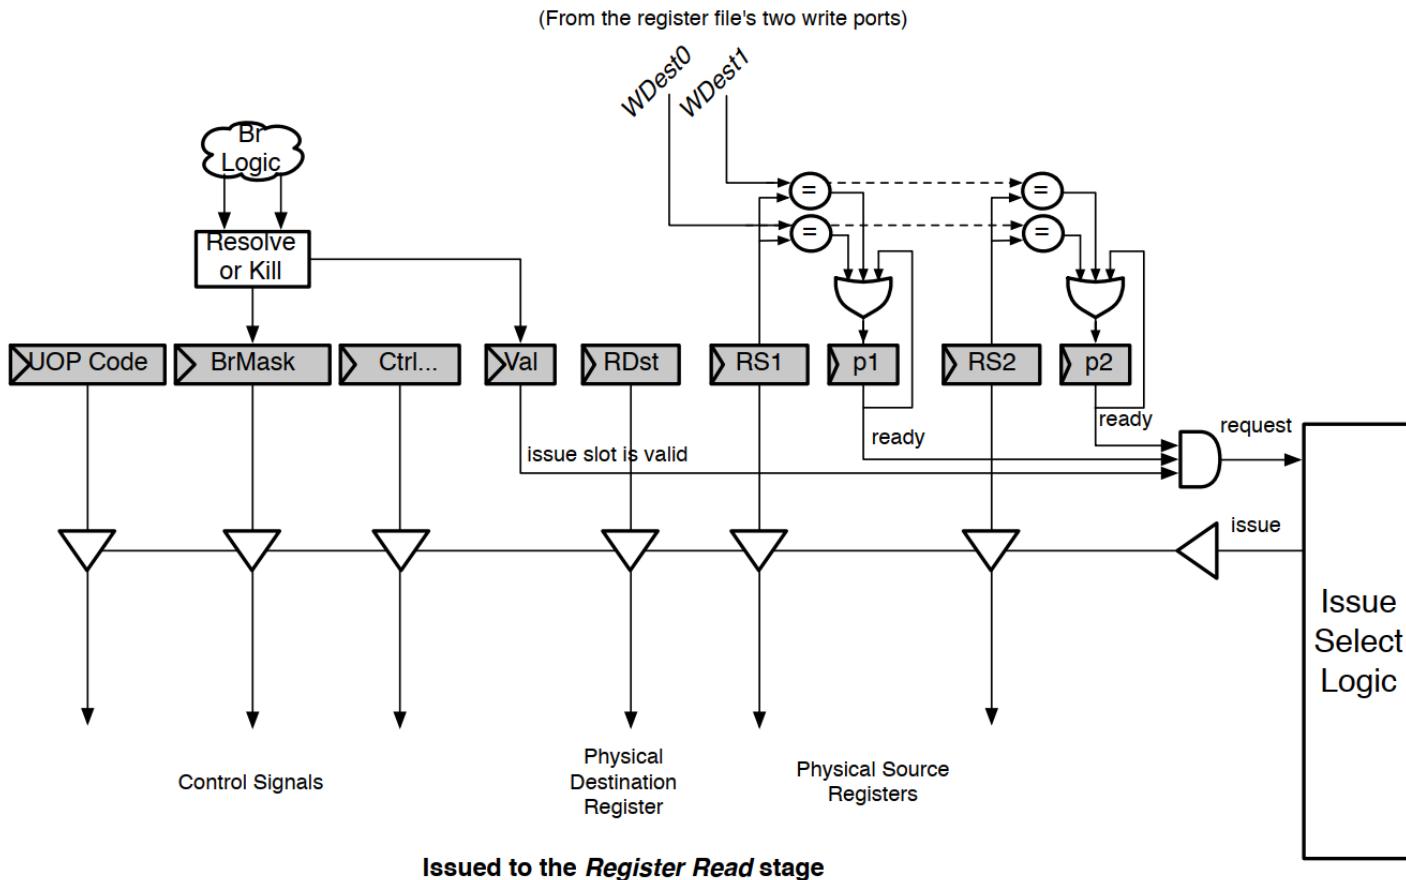
## Register renaming (two parallel channels)



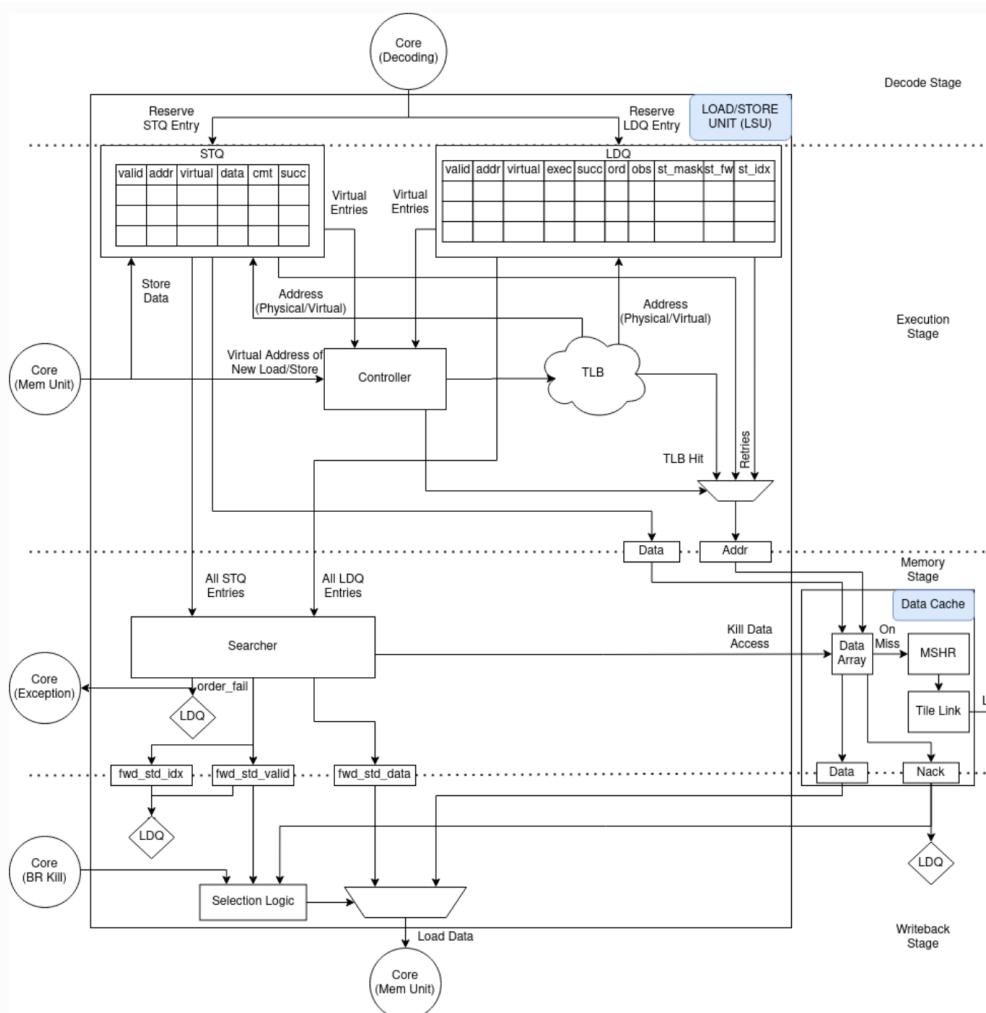
# Dispatching and reorder buffer



## Issue Queue slot



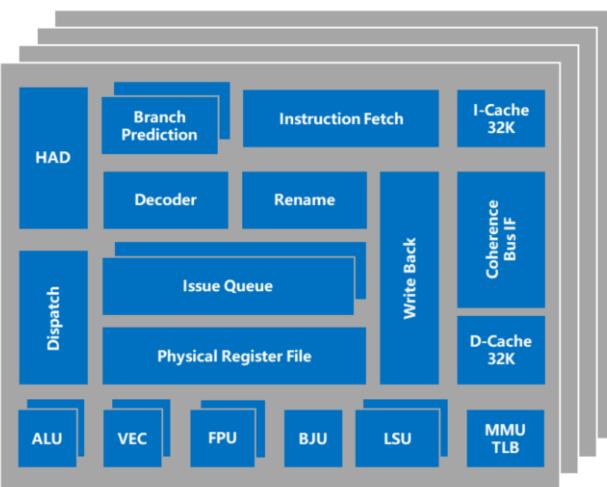
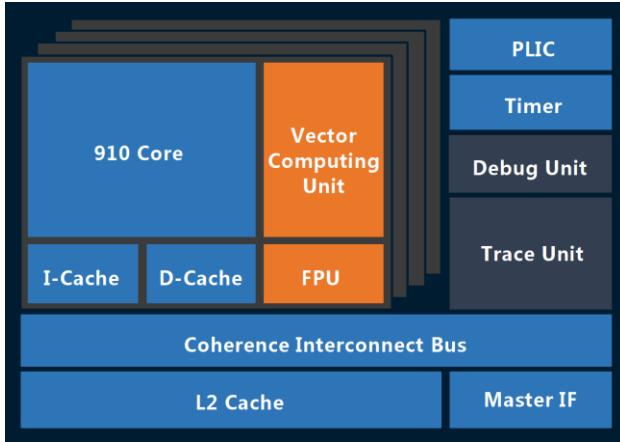
# Load/Store Unit (LSU)





OpenXuantie

# OpenXuantie (Alibaba): design overview



- RISC-V (RV64GCV) architecture
- 4 CPU cores with various capabilities available:  
*OpenE902, OpenE906, OpenC906, OpenC910*
- 64-bit, 12-Stage, Out-of-Order
- Vector Processor Unit
- 50 Custom Instructions
- 32/64KB L1D/I Cache
- Each cluster has a shared, 8/16-way associated L2 cache, up to 8 MB
- Written in Verilog HDL
- Up to 2.5 GHz on TSMC 12 nm
- Fastest open core on release (7,1 CoreMark/MHz)

*C. Chen et al., “Xuantie-910: A Commercial Multi-Core 12-Stage Pipeline Out-of-Order 64-bit High Performance RISC-V Processor with Vector Extension : Industrial Product,” Proc. - Int. Symp. Comput. Archit., vol. 2020-May, pp. 52–64, 2020*

*Presentation:* <https://ieeexplore.ieee.org/document/9220630>

*Repo:* <https://github.com/T-head-Semi>

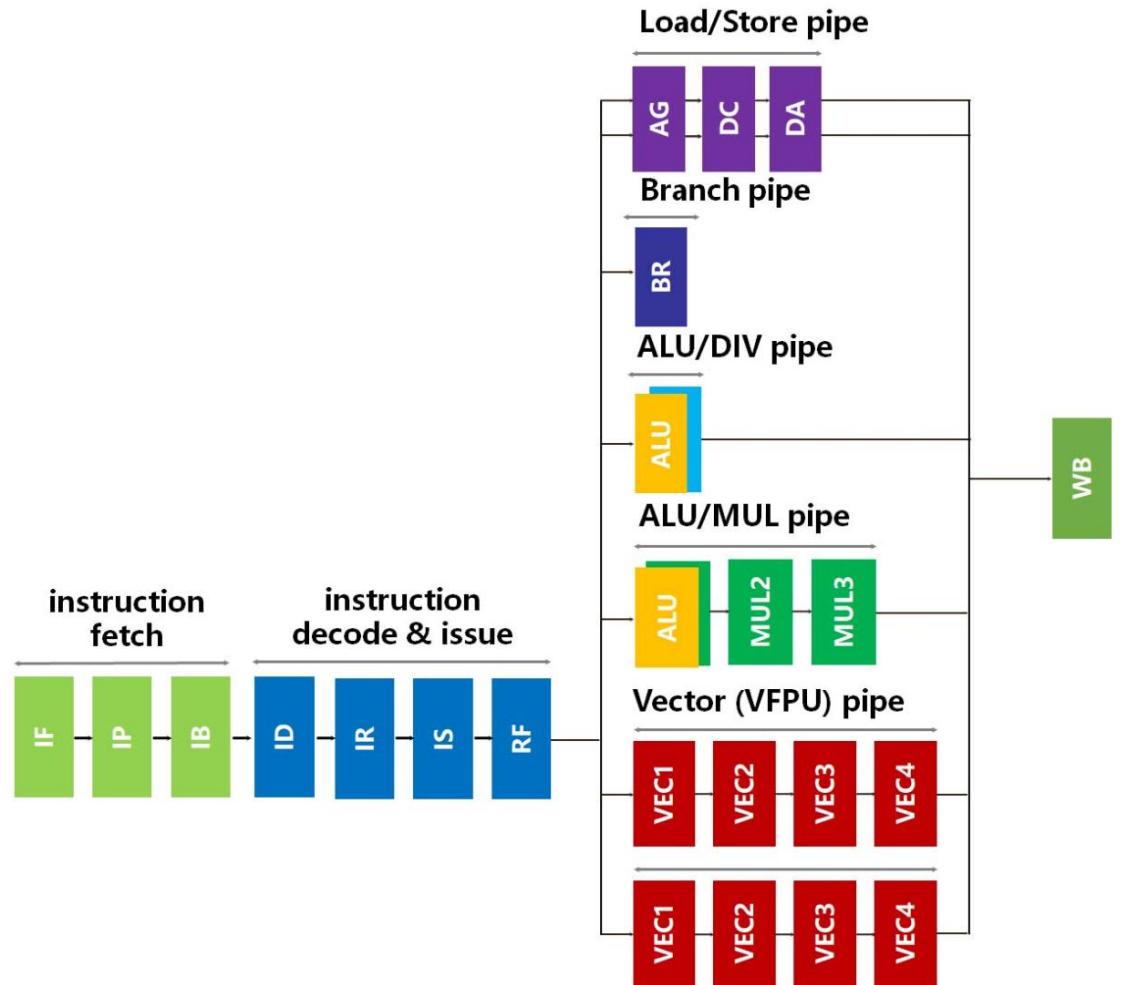
# Pipeline structure

## Frontend:

- Consists of 7 stages
- IFU uses a hybrid branch predictor
- IDU can decode 3 instructions simultaneously and can rename up to 4 instructions using physical registers
- OoO issue engine can issue up to 8 instructions

## Backend:

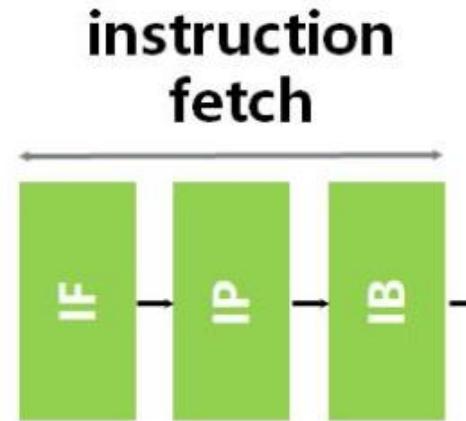
- Multiple execution units (EUs):
  - 2x single-cycle ALUs
  - 1x single cycle branch jump unit
  - 1x dual-issue out-of-order load & store unit
  - 2x FPU and 2x VEC
  - 1x load/store unit (LSU)



# Instruction Fetch Unit (IFU)

## 3 stages:

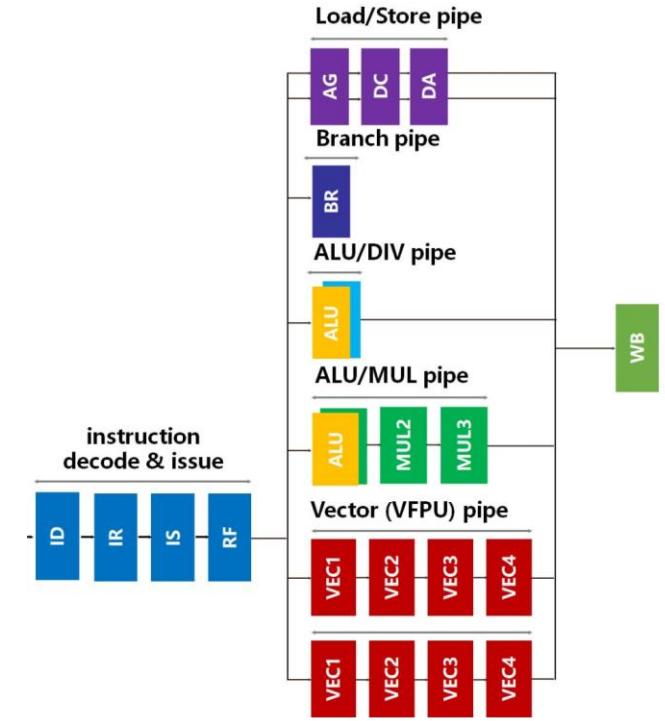
- Instruction Fetch (IF):
  - 128 bit fetch from the L1 cache in a single cycle
- Instruction Pack (IP):
  - pack and pre-process the instructions,
  - process the second-stage of branch jump
  - refill lines in case of cache miss
  - Up to 8 instructions can be packed
  - IDU can maximally decode 3 instructions, the throughput of IP will not become the performance bottleneck
- Instruction Buffer (IB):
  - the packed instruction is buffered
  - up to 3 instructions are sent to the ID pipeline stage
  - the third-stage of jump is processed.



# Execution subsystem

## Execution frontend – 4 stages:

- Instruction Decode (ID):
  - decomposes and decodes instructions
  - splits instruction into micro-instructions
  - zero-latency decoupling of scalar and vector instructions is supported
- Instruction Rename (IR):
  - renames operands in GPR, FGPR and VGPR using physical registers
  - eliminates the use of the costly move instruction
- Instruction Schedule (IS):
  - performs out-of-order instruction scheduling
  - The processor has 8 instruction slots
  - use an Age-Vector based scheduling algorithm
  - support dynamic load balancing by monitoring the workload in the pipeline
- Register File access (RF):
  - reads register file for operands



## Execution backend:

- The EX stage (8 pipes)
- The RT stage:
  - takes care of instruction write-back and retirement
  - The ROB can hold up to 192 instructions

# Memory subsystem: Load/Store Unit (LSU)

- LSU supports:
  - unaligned memory accesses
  - dual Issue OoO memory interface
  - multi-stream prefetching
- Can process one load instruction and one store instruction in parallel
- Requires dedicated load pipe and store pipe
- Each pipe contains 4 pipeline stages
  - address generation (AG)
  - data cache (DC)
  - data alignment (DA)
  - write-back (WB)
- Executions on the load and store pipes access their own RAMs
- Load Queue (LQ) and the Store Queue (SQ) keep the order of instructions

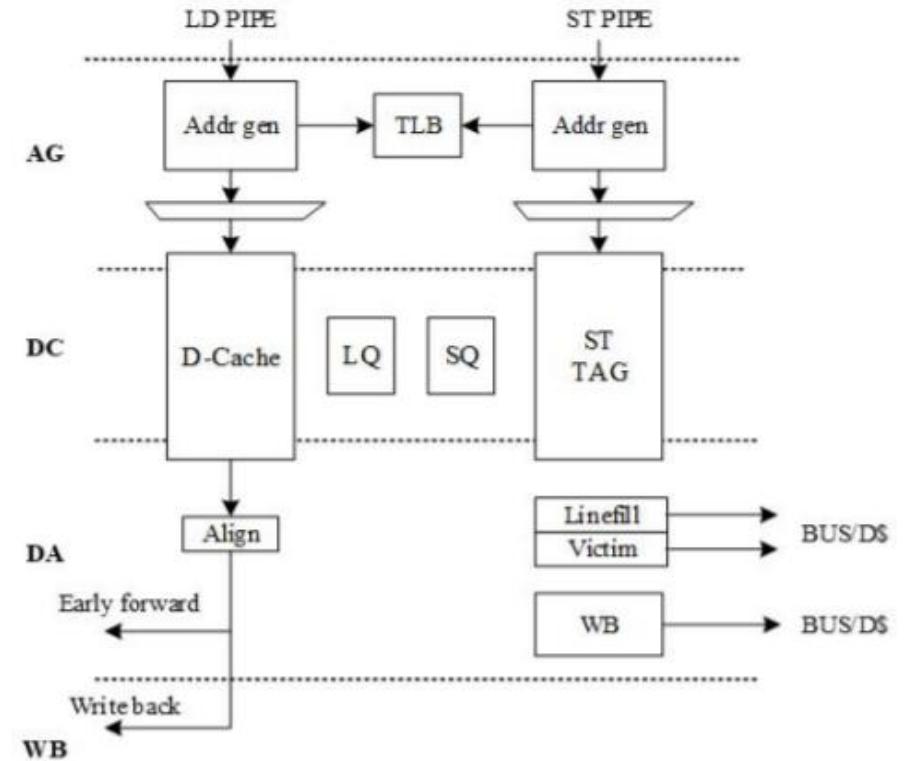
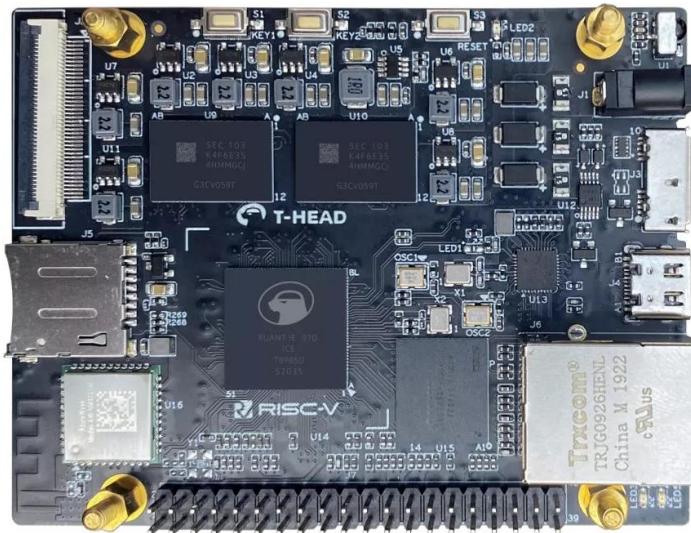


Fig. 9: The pipeline stages in the LSU.

## Xuantie-based products



Lichee RV dev board  
(Sipeed)  
Allwinner D1 SoC  
CPU core: XuanTie C906



RVB-ICE dev board  
(T-Head)  
XuanTie C910 ICE SoC  
CPU core: XuanTie C910

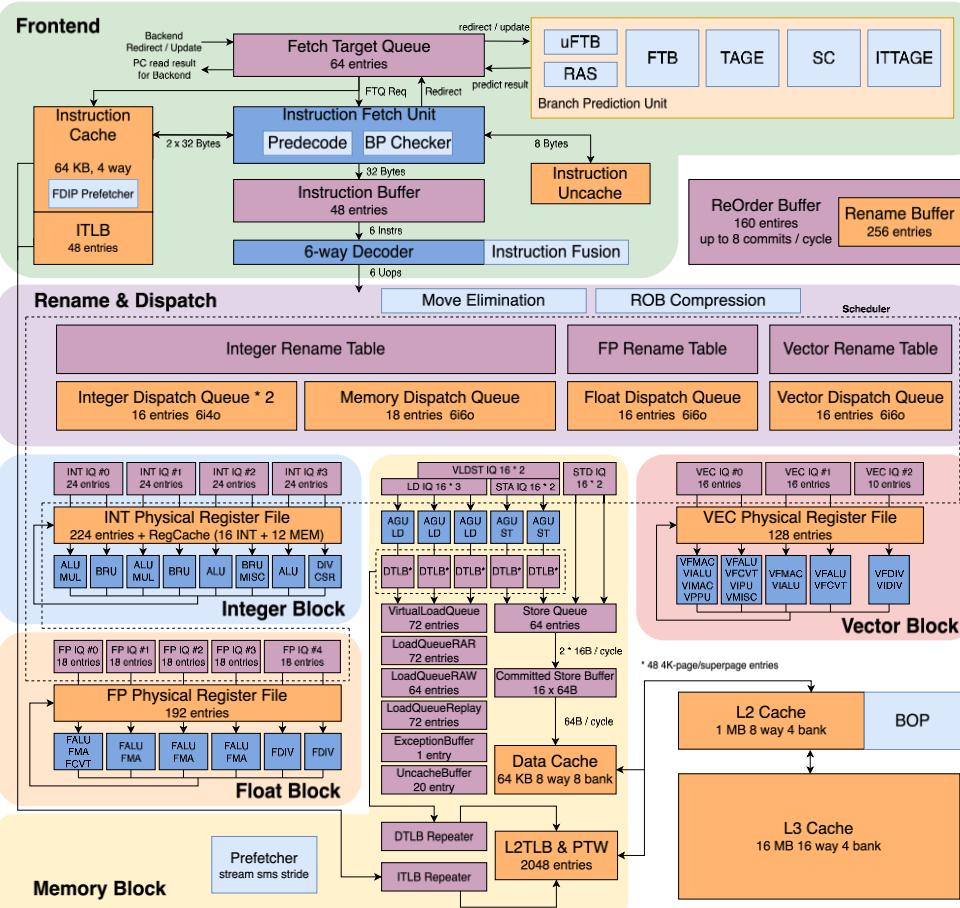


ROMA dev laptop  
(DeepComputing, Xcalibyte)  
Alibaba TH1520 SoC  
CPU core: XuanTie C910



XiangShan

# XiangShan (Chinese Academy of Science): design overview

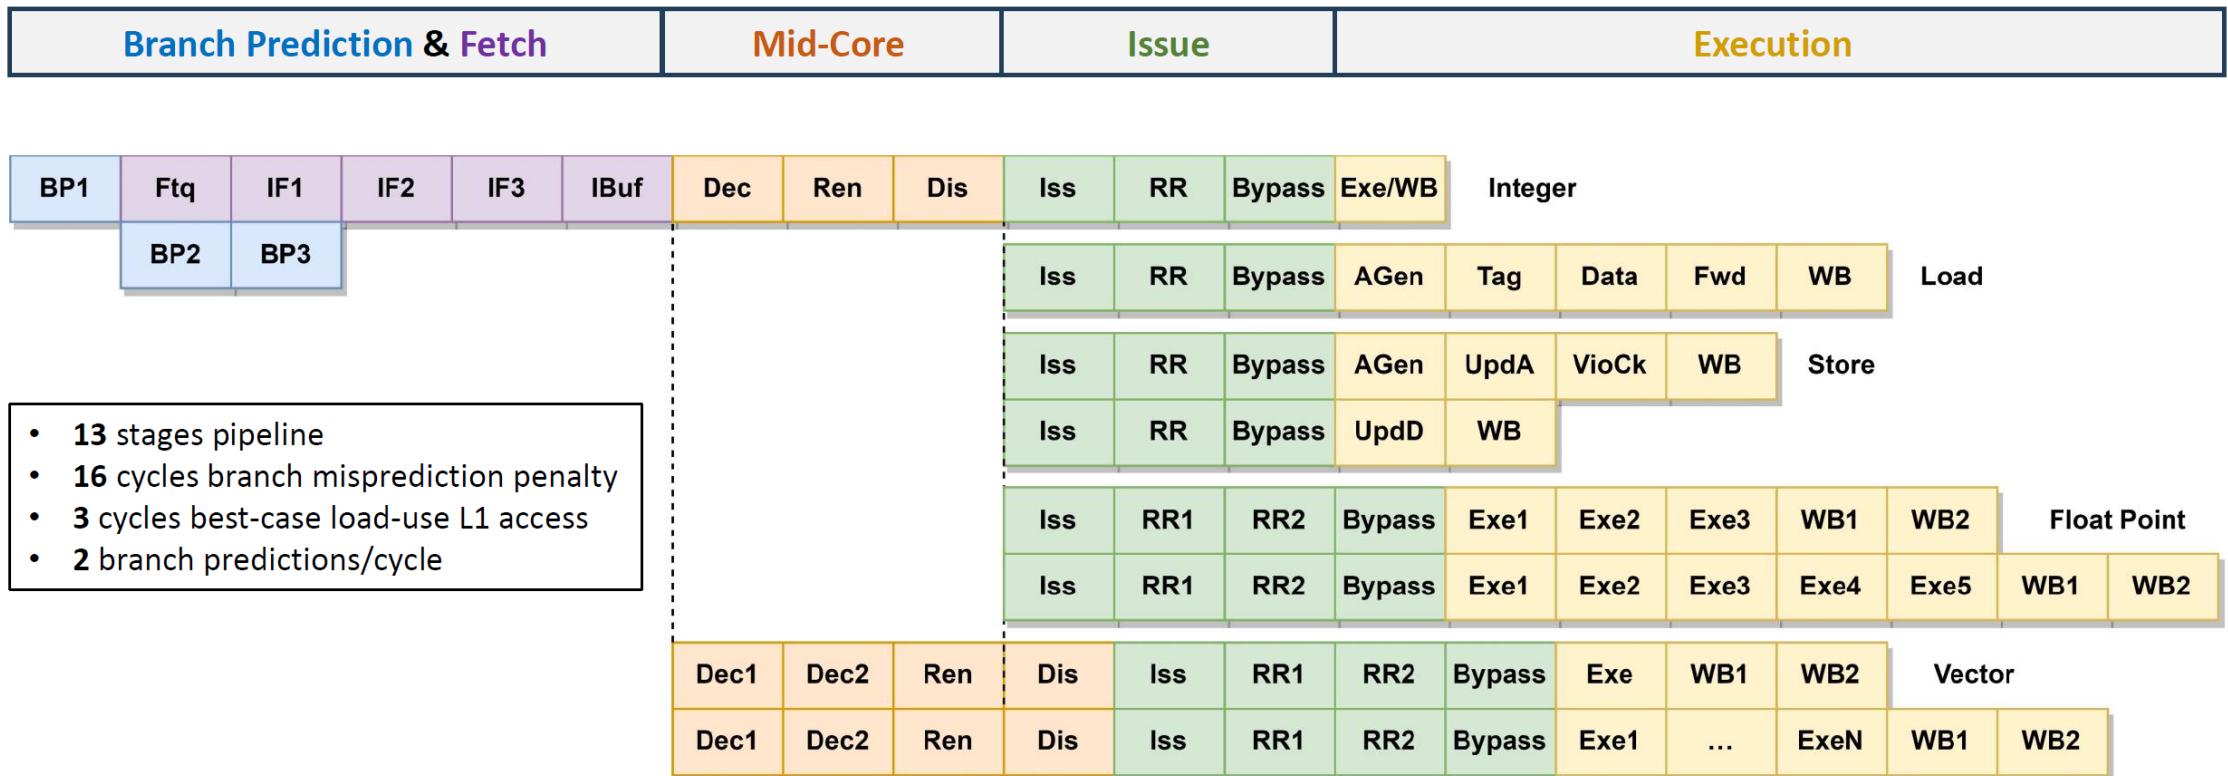


- RISC-V (RV64GCBK) architecture
- Project contains 3x microarchitectures:  
*Yanqihu (stable), Nanhu (stable), Kunminghu (in development)*
- 13-stage 6-issue OoO pipeline
- Boots Debian Linux
- Award-winning project by RISC-V International
- >4100 stars and >550 forks on GitHub
- Written in Chisel HCL, uses components from BOOM
- Had several tape-outs:
  - Yanqihu: TSMC 28nm, 1.3GHz (2022)
  - Nanhu: SMIC(?) 14nm, 2GHz (2023)
- Fastest open core on release (Nanhu: 7,81 CoreMark/MHz)
- Published under Mulan Permissive Software Licence 2 (MulanPSL2)

Documentation: <https://xiangshan.cc>

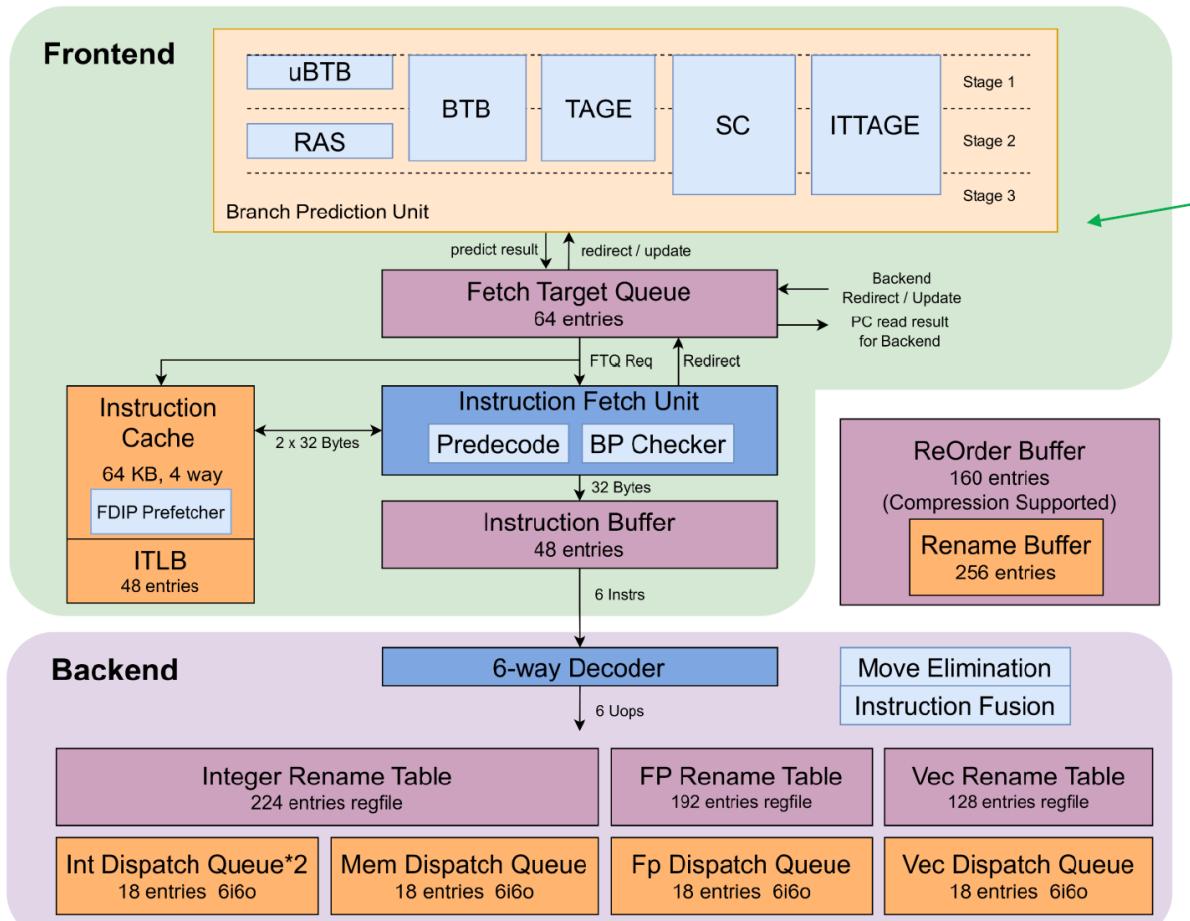
Repo: <https://github.com/OpenXiangShan/XiangShan>

# Pipeline structure (Kunminghu)



Kaifan Wang et al. XiangShan: An Open-Source Project for High-Performance RISC-V Processors Meeting Industrial-Grade Standards. Hot Chips 2024.

# Frontend

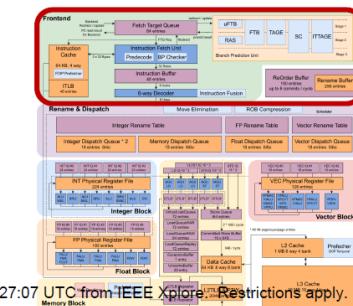


## Multi-level branch predictors

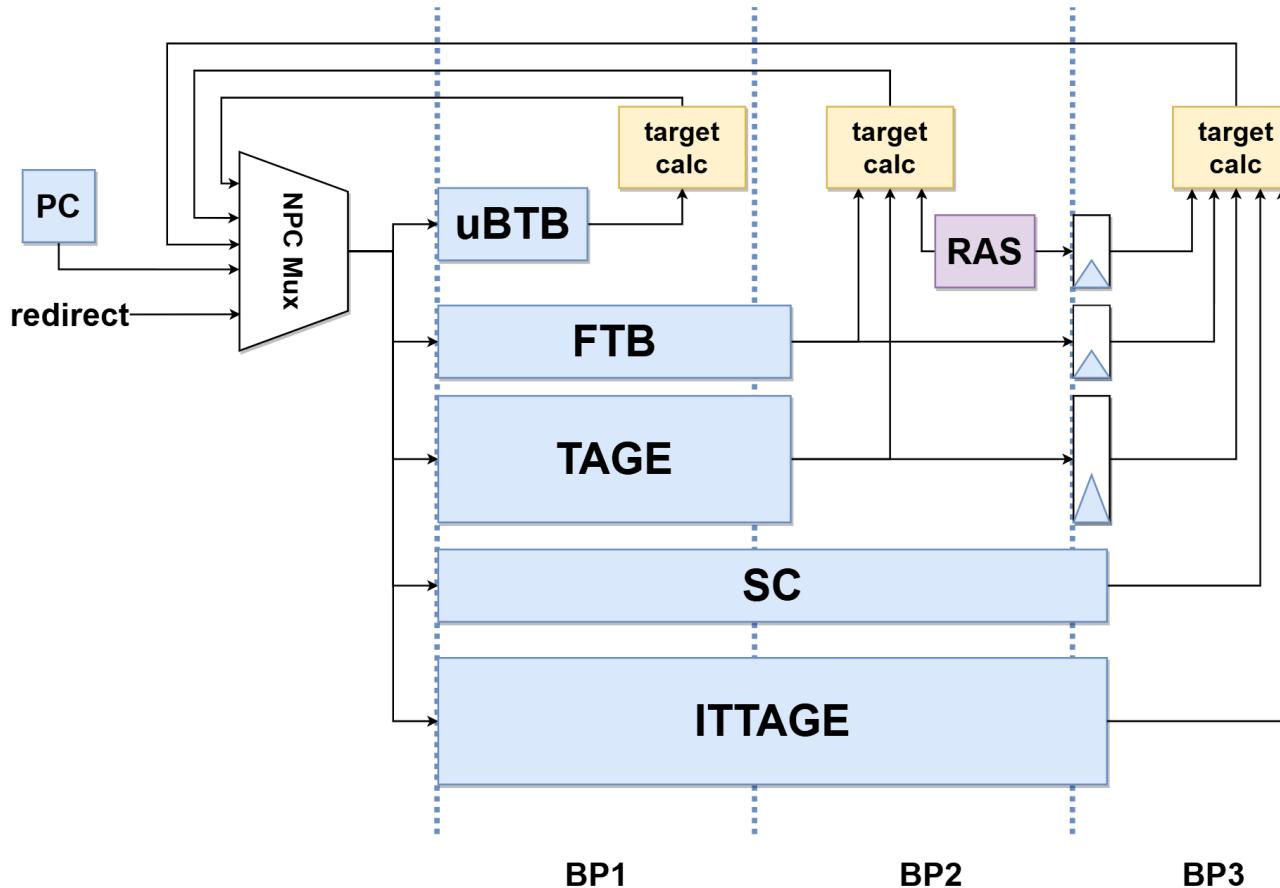
- 256-entry  $\mu$ BTB
- 2K-entry BTB, optional L2 BTB
- 16K TAGE-SC direction predictor
- 2K ITTAGE indirection predictor
- 48-entry return address stack

## Instruction cache and TLB

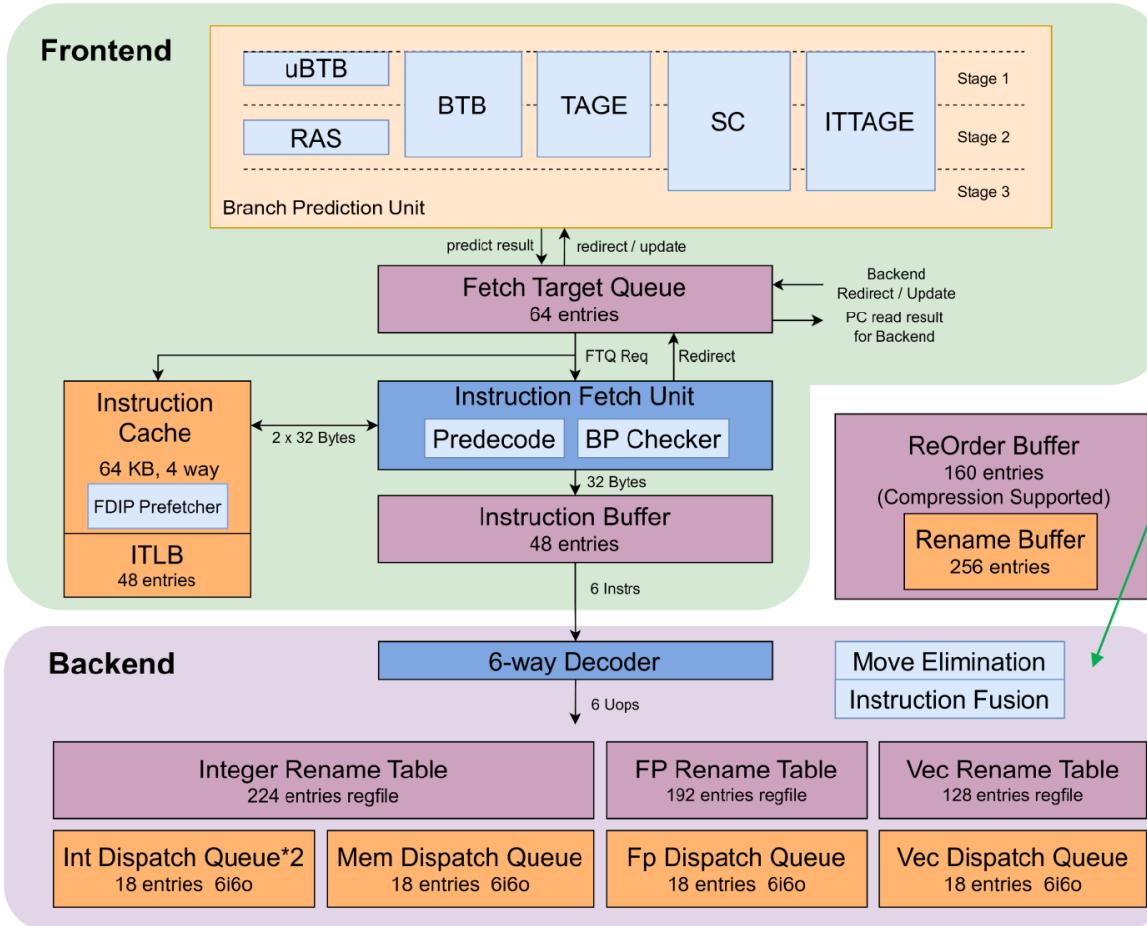
- 64KB 4-way ICache
- 48-entry ITLB
- Fetch-directed instruction prefetcher



## Branch Prediction Unit



# Backend



## 6-wide decode/rename/dispatch

### Register rename

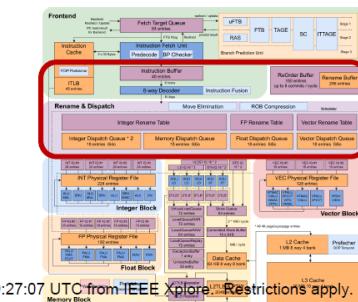
- 224-entry Int regfile
- 192-entry FP regfile
- 128-entry Vec regfile
- Move elimination
- Instruction fusion

### 160 entry ROB

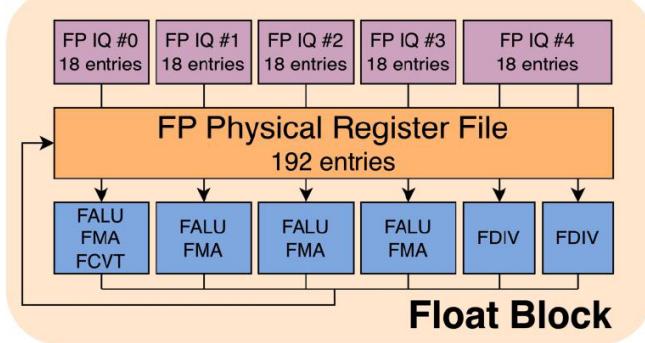
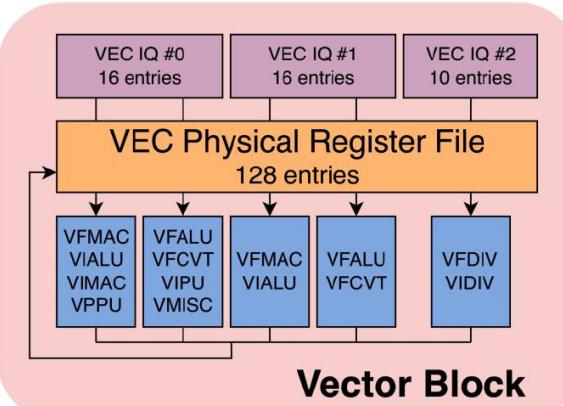
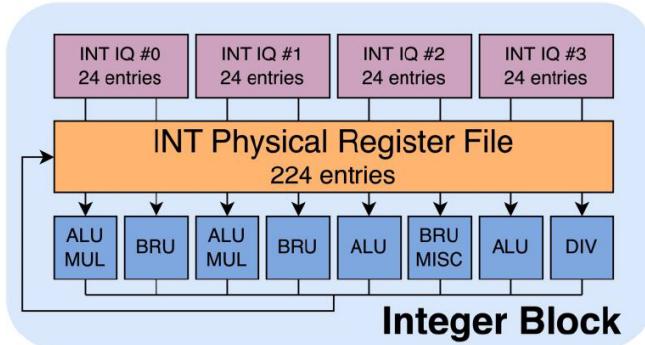
- Support compression (up to 6- $\mu$ op/entry)
- 8-entry retire/cycle
- Recovery via checkpoint + rollback

### 256 entry Rename Buffer

- Bridge the gap between commit & rename table update



# Execution units



## Integer block

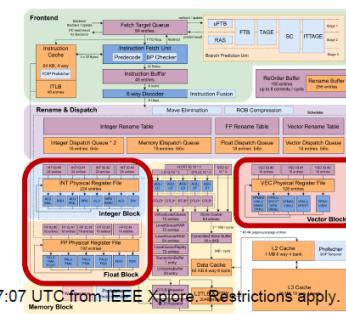
- 4 ALU (2 with 3-stage multiplier)
- 3 BRU for branch resolution
- 1 SRT radix-16 divider

## Float-point block

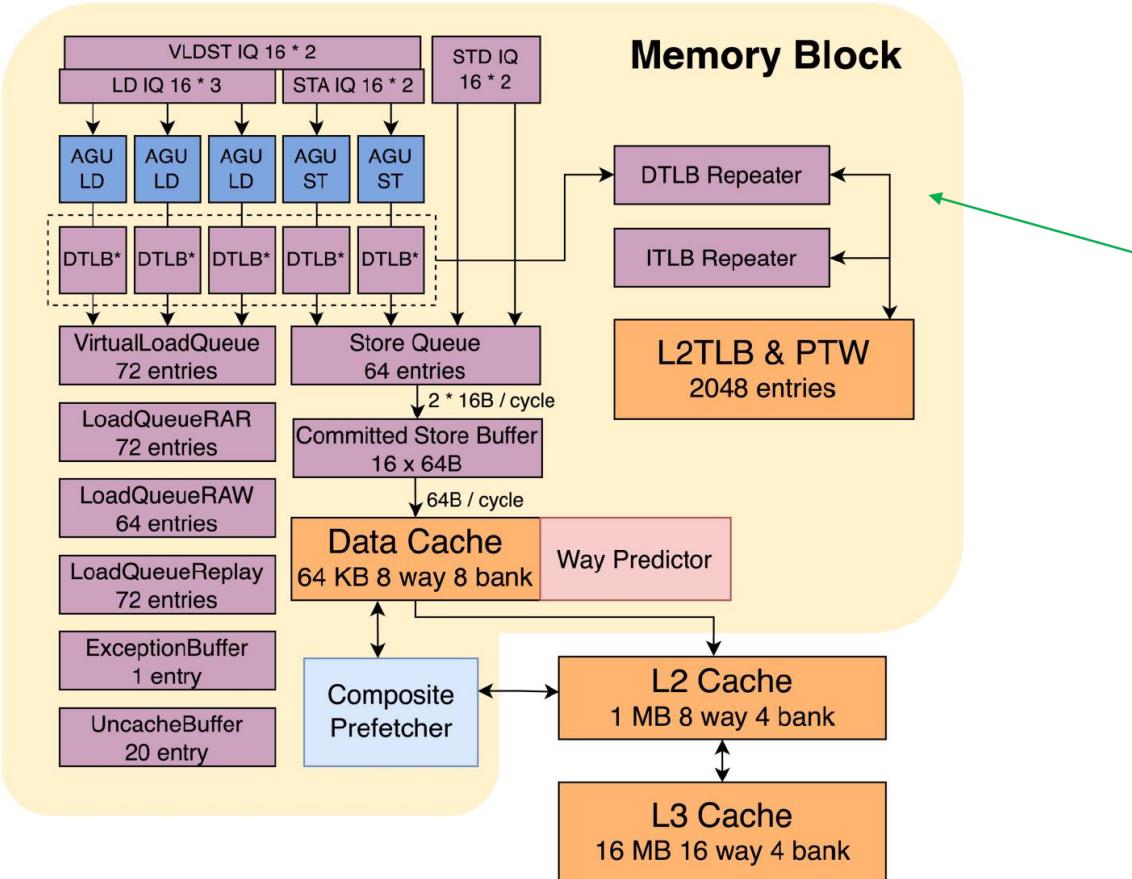
- 4 FPU + 2 FP Div

## Vector block

- 4 VPU + 1 Vec Div
- V1.0 SPEC (VLEN = 128)



# Load/store unit



## Load/Store pipeline & queue

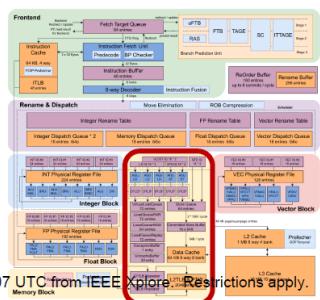
- 3 load pipes, 2 store pipes
- 72 inflight load, 64 inflight store

## MMU

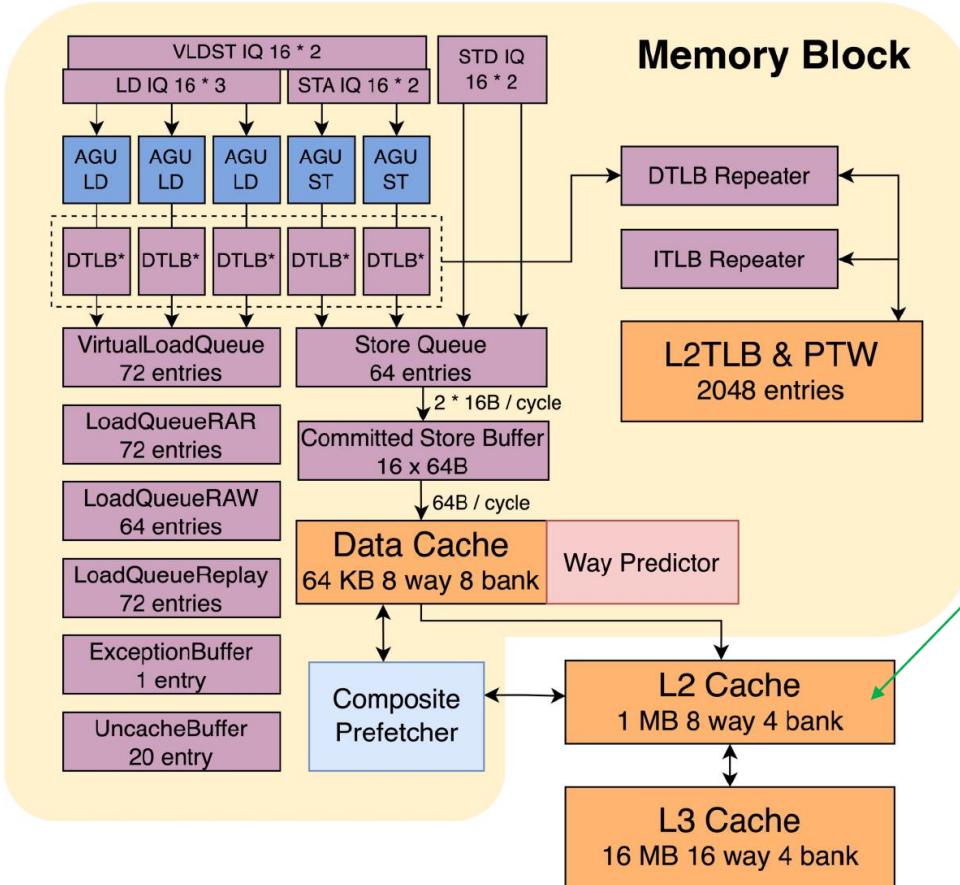
- 48b virtual/52b physical addressing
- 48-entry L1 DTLB, 2K-entry L2 TLB
- Up to 7 outstanding page table walks
- Nested walk for virtualization

## Data cache

- 64KB 8-way VIPT (HW anti-aliasing)
- Way predictor for power efficiency
- Composite prefetcher
  - Stream & Stride prefetching
  - SMS & BOP prefetching
  - Temporal prefetching



# Caches

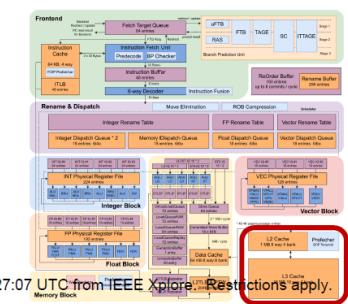


## Private L2 cache

- 8-way, up to **1MB** per core
- Inclusive to D\$, non-inclusive to I\$
- **64+ OTs** (configurable)
- **10-cycle** load-use latency
- PLRU/RRIP replacement

## Shared L3 cache

- 16-way, up to **16MB**
- Inclusive/Non-inclusive to L2\$
- **150+ OTs** (configurable)
- **~30-cycle** load-use latency
- PLRU/RRIP replacement



## Other open-source OoO CPU projects

- riscy-OOO (RISC-V, MIT, in Bluespec HDL)  
paper: <https://people.csail.mit.edu/szzhang/paper/micro2018.pdf>  
repo: <https://github.com/csail-csg/riscy-OOO>
- SweRV EH1 (RISC-V, limited OoO, Western Digital, in SystemVerilog HDL)  
repo: <https://github.com/chipsalliance/Cores-SweRV>
- RSD (RISC-V, Japan, in SystemVerilog HDL)  
repo: <https://github.com/rsd-devel/rsd>



**Thank you for the lesson!**

Alexander Antonov, Assoc. Prof., [antonov@itmo.ru](mailto:antonov@itmo.ru)

Hangzhou, 2025